

CS536_HW1_LeNet_by_cl1288_LIN_chihui

February 29, 2020

1 CS536 Homework Assignment #1

In this homework, we will implement a few core CNN blocks and practice training neural networks by following the guidance step by step. The model to implement is similar to the [LeNet](#). To do this some implementation sketch will be provided on which you can fill in your implementation.

In the following, we first import the basic packages. Feel free to add other packages if necessary. **Note:** The only allowed deep learning framework is PyTorch. Please use Python 3.6 or newer versions and PyTorch 1.3 or newer versions for this homework.

```
[2]: import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader, Dataset
import matplotlib.pyplot as plt
import torchvision.utils as vutils
import numpy as np
```

The following is a sketch of the LeNet class which you will be filling in step-by-step. For now, you don't need to do anything with the following code.

```
[2]: # LeNet sketch code
class LeNet(nn.Module):
    def __init__(self):
        super(LeNet, self).__init__()
        """
        self.c1 = Conv2D()
        self.p2 = nn.AvgPool(2, stride=2)
        self.c3 = Conv2D()
        self.p4 = nn.AvgPool(2, stride=2)
        self.c5 = Linear()
        self.f6 = Linear()
        """

    def forward(self, imgs, labels):
        """
        scores = self.net(imgs)
```

```

o = softmax(scores)
loss = objective(o, labels)
"""
return loss

```

1.1 Task 1: Implement Convolutional Layer

The first task is to implement a convolutional layer by completing the following Conv2D class. The class takes the number of input channels, the number of output channels, stride size, and padding values as inputs.

To-do: - (10 points) Implement code in-between (### start your code here ### and ### End of the code ###)

```

[3]: class Conv2D(nn.Module):
    def __init__(self, dim_in, dim_out, kernel_size, stride, padding, device):
        super(Conv2D, self).__init__()
        """
        inputs:
            dim_in: integer, number of channels in the input
            dim_out: integer, number of channels produced by the convolution
            kernel_size: integer list of length 2, spatial size of the convolving
↪ kernel
            stride: integer list of length 2, stride of the convolution along the the
↪ height dimension and width dimension
            padding: integers list of length 4, zero-padding added to both sides of
↪ the height dimension and width dimension

        """
        # initialize kernel and bias
        self.kernel = nn.Parameter(torch.randn([dim_out, dim_in]+kernel_size,
↪ dtype=torch.float32, device=device)*0.1, requires_grad=True)
        self.bias = nn.Parameter(torch.zeros([dim_out], dtype=torch.float32,
↪ device=device), requires_grad=True)

        self.dim_in = dim_in
        self.dim_out = dim_out

        self.kernel_size = kernel_size
        self.stride = stride
        self.padding = padding
        self.device = device

    def conv2d_forward(self, X):
        """
        inputs:

```

```

    X: input images
    outputs:
    Y: output produced by the convolution
    """

    ### Star your code here ###
    # Padding
    pad_x = torch.zeros( X.shape[0], X.shape[1],
                        X.shape[2]+self.padding[0]+self.padding[1],
                        X.shape[3]+self.padding[2]+self.padding[3], dtype=X.
↪dtype)
    pad_x[:, :, self.padding[0]:X.shape[2]+self.padding[0], self.padding[2]:X.
↪shape[3]+self.padding[2]] = X
    # convolutional computation: (padding input)* (the kernel)
    pad_x_unfold = pad_x.unfold(2, self.kernel_size[0], self.stride[0]).
↪unfold(3, self.kernel_size[1], self.stride[1]).unsqueeze(1)
    kernel_unsqueeze = self.kernel.unsqueeze(0).unsqueeze(3).unsqueeze(4)
    x_kernel = (pad_x_unfold.cuda() * kernel_unsqueeze.cuda()).sum(-1).sum(-1).
↪sum(2).cuda()

    Y = (x_kernel.cuda()+self.bias.unsqueeze(0).unsqueeze(2).unsqueeze(3).
↪expand_as(x_kernel).cuda())
    ### End of the code ###
    return Y

def forward(self, x):
    return self.conv2d_forward(x)

```

1.2 Conv2D Correctness Check

Run the correctness checking code. If your implementation is correct, you should be able to see the output as follows:

```

tensor([[[[ 1.0519,  1.3811],
           [ 2.5701,  2.0508]],

        [[ 1.3159,  2.4203],
           [ 8.4296, 10.1662]]],

       [[[ 4.4640,  5.9235],
           [ 4.5179,  2.3144]],

        [[13.1246, 15.5932],
           [20.2635, 22.9700]]]], grad_fn=<AddBackward0>)

```

```
[4]: # correctness checking
torch.random.manual_seed(0)
x = torch.arange(50).view(2,1,5,5).float()
my_conv = Conv2D(1,2,[3,3],[3,3],[1,1,1,1], torch.device('cpu'))
y = my_conv(x)
print(y)

tensor([[[[ 1.0519,  1.3811],
           [ 2.5701,  2.0508]],

         [[ 1.3159,  2.4203],
           [ 8.4296, 10.1662]]],

        [[[ 4.4640,  5.9235],
           [ 4.5179,  2.3144]],

         [[13.1246, 15.5932],
           [20.2635, 22.9700]]]], device='cuda:0', grad_fn=<AddBackward0>)
```

1.3 Guide: Pooling Layer

We will not implement the pooling layer. Instead, we will use the Pytorch API (`torch.nn.AvgPool2d`). Feel free to implement it by yourself (if you want). For the detail information about the pooling API, check the [documents](#).

1.4 Task 2: Implementing Linear Layer

Complete the following linear layer module. To specify a linear layer, input dimension and output dimension are provided. The linear layer performs the following computation: $y = xW^T + b$.

To-do: - (5 points) Implement code in-between (### start your code here ### and ### End of the code ###)

```
[4]: class Linear(nn.Module):
    def __init__(self, dim_in, dim_out, device):
        super(Linear, self).__init__()

        self.weights = nn.Parameter(torch.randn([dim_out, dim_in], dtype=torch.
↪float32, device=device)*0.1, requires_grad=True)

        self.bias = nn.Parameter(torch.zeros([dim_out], dtype=torch.float32,
↪device=device), requires_grad=True)
        self.dim_out = dim_out
        self.device = device
```

```

def linear_forward(self, X):
    """
    inputs:
        X: tensor of shape (batch_size, *, dim_in)
    outputs:
        Y: tensor of shape (batch_size, *, dim_out)
    """

    ### Star your code here ###
    Y = torch.mm(X, self.weights.t()) + self.bias
    ### End of the code ###
    return Y

def forward(self, X):
    return self.linear_forward(X)

```

1.5 Linear Correctness Check

Run the following correctness checking code. If your implementation is correct, you should be able to see the output as follows:

```

tensor([[ -2.1595,  -0.2037,   1.8567],
        [ -6.9537,   0.7306,   2.5298],
        [-11.7479,   1.6648,   3.2028],
        [-16.5422,   2.5991,   3.8759],
        [-21.3364,   3.5334,   4.5489]], grad_fn=<AddBackward0>)

```

```

[6]: # correctness checking
torch.random.manual_seed(0)
x = torch.arange(50).view(5, 10).float()
my_linear = Linear(10, 3, torch.device('cpu'))
y = my_linear(x)
print(y)

```

```

tensor([[ -2.1595,  -0.2037,   1.8567],
        [ -6.9537,   0.7306,   2.5298],
        [-11.7479,   1.6648,   3.2028],
        [-16.5422,   2.5991,   3.8759],
        [-21.3364,   3.5334,   4.5489]], grad_fn=<AddBackward0>)

```

1.6 Task 3: Loss Functions and SGD

The loss function for classification task is the Cross-Entropy Loss. For this, we need to implement the softmax output layer first and then the cross-entropy loss.

Softmax function normalizes the output so that its sum becomes 1 and each output is nonnegative:

$$\hat{\mathbf{y}} = \text{softmax}(\mathbf{o}), \text{ where } \hat{y}_i = \frac{\exp(o_i)}{\sum_j \exp(o_j)}$$

Cross-Entropy Loss is as the objective function for this classification task. When the loss is minimized, the likelihood function will be maximized:

$$l = -\log P(y | x) = -\sum_j y_j \log \hat{y}_j$$

To-do: - (10 points) Complete the function `softmax1d()`. - (10 points) Complete the function `cross_entropy_loss()`.

Note that you should implement the function using primitive PyTorch APIs such as `exp()` and `matmul()`, instead of simply using pythor API for softmax and `cross_entropy_loss`.

```
[5]: def softmax1d(scores):
    """
    inputs:
        scores: (N, C), predicted scores for each input, where N is the number of
        ↪ samples and C is the number of
            classes.
    outputs:
        p: (N, C), probability distribution over classes. Converted from input
        ↪ (scores) with a softmax operation.

    Note: Do be careful of the numerical error!
    """
    ### Star your code here ###
    scores = scores - scores.max(1)[0].unsqueeze(1).expand_as(scores)
    exp_scores = torch.exp(scores)
    p = exp_scores / (exp_scores.sum(1).unsqueeze(1))
    ### End of the code ###

    return p

def cross_entropy_loss(pred_score, labels):
    """
    inputs:
        pred_score: (N, C), probaility distribution or pred_scores over classes,
        ↪ where N is the number of samples and C is the number of
            classes.
    outputs:
        loss: (N,), cross entropy loss for each sample.

    Note: Do be careful of the numerical error!
    """
    ### Star your code here ###
```

```

    label_onehot = torch.cuda.FloatTensor(labels.shape[0], 10).zero_().
    ↪scatter_(1, labels.unsqueeze(1).cuda(), 1).cuda()
    loss = -1*(label_onehot.cuda()*torch.log(pred_score+ 1e-9).cuda()).sum(1)
    ### End of the code ###

    return loss

```

To-do: - (10 points) Next task is to implement the update rule of stochastic gradient descent. Complete the following function.

```

[6]: def step(weights, grad, lr):
    """
    inputs:
    weights: list of learnable parameters
    grad: list of gradient of the loss w.r.t the learnable parameters
    lr: learning rate for gradient descent
    outputs:
    None. Make sure updating the weights with in-place operation, e.g. tensor.
    ↪add_(). No output need be returned.
    """
    ### Star your code here ###
    for i in range(len(weights)):
        weights[i].data.add_(-grad[i].data.clone().detach()*lr)
    ### End of the code ###

```

1.7 Task 4: LeNet Forward Pass

Using the above components required to implement the LeNet, we can complete the LeNet class as follows.

To-do: - (20 points) Complete the function **forward()** which takes the input images and labels and outputs the cross-entropy loss (for the batch) and predicted distribution. For more details, refer to the comments below.

```

[9]: # LeNet sketch code
class LeNet(nn.Module):
    def __init__(self, img_c, device):
        super(LeNet, self).__init__()
        self.c1 = Conv2D(img_c, 6, [5,5], [1,1], [2,2,2,2], device)
        self.p2 = nn.MaxPool2d(2, stride=2)
        self.c3 = Conv2D(6, 16, [5,5], [1,1], [0,0,0,0], device)
        self.p4 = nn.MaxPool2d(2, stride=2)
        self.f5 = Linear(400, 120, device)
        self.f6 = Linear(120, 84, device)
        self.f7 = Linear(84, 10, device)
        self.device = device

```

```

def forward(self, imgs, labels):
    """
    inputs:
        imgs: (N, C, H, W), training samples from the MNIST training set, where N
        → is the number of samples (batch_size),
        C is the image color channle number, H and W are the spatial size of
        → the input images.
        labels: (N, L), ground truth for the input images, where N is the number
        → of samples (batch_size) and L is the
        number of classes.
    outputs:
        loss: (1,), mean loss value over this batch of inputs.

    """
    N = imgs.shape[0]

    o_c1 = F.relu(self.c1(imgs))
    o_p2 = self.p2(o_c1)
    o_c3 = F.relu(self.c3(o_p2))
    o_p4 = self.p4(o_c3)

    ### Start the code here ###
    # 1. Please complete the rest of LeNet to get the scores predicted by LeNet
    → for each input images #
    # need to flatten the matrix before forwarding to the dense layer
    o_f5 = F.relu(self.f5(o_p4.reshape(o_p4.shape[0], o_p4.shape[1]*o_p4.
    → shape[2]*o_p4.shape[3])))
    o_f6 = F.relu(self.f6(o_f5))
    o_f7 = self.f7(o_f6)

    # 2. Please use the implemented objective function to obtain the losses of
    → each input. #
    p = softmax1d(o_f7)

    # 3. We will return the mean value of the losses. #
    loss = cross_entropy_loss(p, labels)

    ### End of the code ###
    return loss.mean(), p

```


1.8 Guide: Dataset Preparation

We use MNIST dataset to train the LeNet. Run the following cell to get the dataset ready for the training. Change the `data_path` to a proper one.

```
[12]: import os
import urllib.request

data_path = './CS536_MNIST/'
if not os.path.exists(data_path):
    os.mkdir(data_path)
    print("Starting downloading MNIST to {}".format(data_path))

import urllib
dataset_dict = {
    'train_images': "http://yann.lecun.com/exdb/mnist/
→train-images-idx3-ubyte.gz",
    'train_labels': "http://yann.lecun.com/exdb/mnist/
→train-labels-idx1-ubyte.gz",
    'test_images': "http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.
→gz",
    'test_labels': "http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.
→gz",}

for f, url in dataset_dict.items():
    urllib.request.urlretrieve(url, data_path + f)
```

The class `dataset` has been provided below. More about PyTorch Dataset can be found [here](#). Please run the following Jupyter cell to make sure the dataset is ready for training.

```
[13]: train_img_file = data_path + 'train_images'
train_lb_file = data_path + 'train_labels'
test_img_file = data_path + 'test_images'
test_lb_file = data_path + 'test_labels'

class MNISTDataset(Dataset):
    def __init__(self, ds_size=10000, split='training'):
        self.split = split
        if self.split == 'training':
            img_file = train_img_file
            lb_file = train_lb_file
            n_samples = 60000
        else:
            img_file = test_img_file
            lb_file = test_lb_file
            n_samples = 10000
        self.ds_size = ds_size
```

```

import gzip
with gzip.open(img_file, 'rb') as f:
    imgs = f.read()
imgs = np.frombuffer(imgs[16:], dtype=np.uint8).astype(np.float32)
with gzip.open(lb_file, 'rb') as f:
    lb = f.read()
lbs = np.frombuffer(lb[8:], dtype=np.uint8).astype(np.float32)

imgs = torch.tensor(imgs).view(n_samples, 1, 28, 28) - 125.
lbs = torch.tensor(lbs).long()

self.imgs = imgs[:ds_size]
self.lbs = lbs[:ds_size]

def __len__(self):
    return self.ds_size

def __getitem__(self, idx):
    return self.imgs[idx], self.lbs[idx]

```

1.9 Guide: Training for Overfitting

First, we will make our model overfit. It is a good practice to check if a model can overfit well (It should do it well in a proper setting. If not, your model may have some bug, the model complexity is too simple, or some training parameters like learning rate are not good.)

For this, we will first use only 1000 data points instead of the full dataset (smaller datasets makes models overfit more easily if the model complexity is fixed.)

We also use the 20% -split as the validation set. The dataset loading script is provided below.

```

[12]: dataset_size = 1000
validation_size = int(0.2 * 1000)

ds = MNISTDataset(ds_size=dataset_size)

# split the dataset into training set and validation set
train_ds, val_ds = torch.utils.data.random_split(ds, [dataset_size -
    ↪ validation_size, validation_size])

# training batch size, hyper-parameter
batch_size = 24

# dataset loader

```

```

train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
    ↪drop_last=True)
val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=True, drop_last=True)
img_channel = 1

```

The training script is provided as follows. After 100 epochs, you will see LeNet is overfitted as the validation error is a lot larger than the training error. Also, the classification accuracy is almost 100% on the training set, while only around 70% on the validation set.

```

[13]: # learning rate, hyper-parameter
lr = 1e-3

# using GPU if it's available
# device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device = torch.device('cuda')

net = LeNet(img_channel, device)
# keep a list of model parameters
params = [p for (n, p) in net.named_parameters()]

# training epochs, hyper-parameter
epochs = 100

# keep tracking of the changing of loss and accuracy of predictions
train_loss_list = []
val_loss_list = []
train_acc_list = []
val_acc_list = []

# the printing frequency, feel free to change this
print_interval = 50
for e in range(epochs):
    net.train()
    for i, (imgs, lbs) in enumerate(train_dl):
        imgs = imgs.to(device)
        lbs = lbs.to(device)
        loss, prob = net(imgs, lbs)

        net.zero_grad()
        grad = torch.autograd.grad(loss, params)

        # update weights
        step(params, grad, lr)

        # obtain the predictions
        pred = prob.argmax(dim=-1).view(batch_size)
        acc = (pred == lbs).float().mean()

```

```

if i % print_interval == 0:
    print("step {}, loss {}".format(i + e*len(train_dl), loss))
    print("Target:\t {} \nPred:\t {}".format(lbs[:8], pred[:8]))
    # visualize some samples
    imgs_to_vis = vutils.make_grid(imgs[:8].cpu()+125., nrow=8, pad_value=1)
    plt.imshow(imgs_to_vis.permute(1,2,0).numpy().astype(np.uint8))
    plt.axis("off")
    plt.show()

train_loss_list.append(loss.detach().mean())
train_acc_list.append(acc.detach().mean())

net.eval()
for i, (imgs, lbs) in enumerate(val_dl):
    imgs = imgs.to(device)
    lbs = lbs.to(device)
    loss, prob = net(imgs, lbs)

    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()

val_loss_list.append(loss.detach().mean())
val_acc_list.append(acc.detach().mean())

# plotting logs
plt.plot(np.arange(epochs), train_loss_list, '-r',
         np.arange(epochs), val_loss_list, '-g')
plt.legend(('training error', 'validation error'))
plt.show()
plt.plot(np.arange(epochs), train_acc_list, '-r',
         np.arange(epochs), val_acc_list, '-g')
plt.legend(('training acc', 'validation acc'))
plt.show()

```

```

step 0, loss 19.859798431396484
Target:  tensor([1, 7, 9, 6, 3, 1, 1, 4], device='cuda:0')
Pred:    tensor([0, 0, 0, 0, 0, 0, 0, 0], device='cuda:0')

```



step 33, loss 17.269390106201172

Target: tensor([9, 0, 8, 1, 4, 6, 7, 2], device='cuda:0')

Pred: tensor([4, 4, 4, 4, 4, 4, 4, 4], device='cuda:0')

A horizontal strip showing a sequence of handwritten digits on a black background. The digits are white and appear to be '9', '0', '8', '1', '4', '6', '7', and '2' from left to right.

step 66, loss 20.722686767578125

Target: tensor([9, 3, 1, 0, 4, 4, 4, 4], device='cuda:0')

Pred: tensor([6, 6, 6, 6, 6, 6, 6, 6], device='cuda:0')

A horizontal strip showing a sequence of handwritten digits on a black background. The digits are white and appear to be '9', '3', '1', '0', '4', '4', '4', and '4' from left to right.

step 99, loss 14.678973197937012

Target: tensor([1, 6, 5, 2, 6, 6, 2, 7], device='cuda:0')

Pred: tensor([3, 6, 6, 0, 6, 6, 0, 3], device='cuda:0')

A horizontal strip showing a sequence of handwritten digits on a black background. The digits are white and appear to be '1', '6', '5', '2', '6', '6', '2', and '7' from left to right.

step 132, loss 10.005305290222168

Target: tensor([6, 8, 5, 8, 6, 0, 0, 3], device='cuda:0')

Pred: tensor([6, 4, 6, 2, 6, 0, 6, 3], device='cuda:0')

6 8 5 8 6 0 0 3

step 165, loss 13.06640911102295

Target: tensor([7, 4, 6, 0, 0, 1, 1, 5], device='cuda:0')

Pred: tensor([4, 4, 2, 0, 0, 2, 2, 2], device='cuda:0')

7 4 6 0 0 1 1 5

step 198, loss 9.484553337097168

Target: tensor([5, 3, 3, 9, 9, 9, 5, 6], device='cuda:0')

Pred: tensor([4, 3, 4, 2, 4, 4, 4, 6], device='cuda:0')

5 3 3 9 9 9 5 6

step 231, loss 12.914787292480469

Target: tensor([4, 5, 8, 7, 8, 1, 4, 7], device='cuda:0')

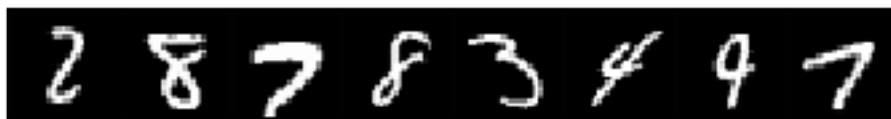
Pred: tensor([4, 4, 6, 3, 3, 1, 4, 4], device='cuda:0')

4 5 8 7 8 1 4 7

step 264, loss 7.584327220916748

Target: tensor([2, 8, 7, 8, 3, 4, 9, 7], device='cuda:0')

Pred: tensor([1, 5, 3, 1, 3, 1, 7, 7], device='cuda:0')

A horizontal strip showing the handwritten digits 2, 8, 7, 8, 3, 4, 9, and 7 in white on a black background.

step 297, loss 6.807391166687012

Target: tensor([7, 7, 6, 1, 2, 1, 4, 4], device='cuda:0')

Pred: tensor([7, 7, 6, 1, 2, 1, 4, 4], device='cuda:0')

A horizontal strip showing the handwritten digits 7, 7, 6, 1, 2, 1, 4, and 4 in white on a black background.

step 330, loss 6.049202919006348

Target: tensor([4, 2, 1, 3, 8, 2, 9, 0], device='cuda:0')

Pred: tensor([4, 2, 1, 3, 2, 2, 4, 0], device='cuda:0')

A horizontal strip showing the handwritten digits 4, 2, 1, 3, 8, 2, 9, and 0 in white on a black background.

step 363, loss 1.3733537197113037

Target: tensor([5, 0, 4, 7, 4, 1, 6, 3], device='cuda:0')

Pred: tensor([5, 0, 4, 4, 4, 1, 6, 3], device='cuda:0')

A horizontal strip showing the handwritten digits 5, 0, 4, 7, 4, 1, 6, and 3 in white on a black background.

step 396, loss 3.6519713401794434

Target: tensor([7, 5, 2, 6, 5, 3, 2, 9], device='cuda:0')

Pred: tensor([7, 5, 2, 6, 5, 3, 2, 7], device='cuda:0')



step 429, loss 0.5592606663703918

Target: tensor([4, 4, 5, 0, 6, 1, 6, 6], device='cuda:0')

Pred: tensor([4, 4, 5, 0, 6, 7, 6, 6], device='cuda:0')



step 462, loss 1.0333976745605469

Target: tensor([2, 1, 7, 2, 8, 5, 8, 6], device='cuda:0')

Pred: tensor([2, 1, 7, 8, 8, 5, 8, 8], device='cuda:0')



step 495, loss 0.21424254775047302

Target: tensor([1, 4, 3, 9, 7, 1, 1, 4], device='cuda:0')

Pred: tensor([1, 4, 3, 4, 7, 1, 1, 4], device='cuda:0')

1 4 3 9 7 1 1 4

step 528, loss 0.08810263127088547

Target: tensor([8, 9, 8, 7, 6, 5, 3, 1], device='cuda:0')

Pred: tensor([8, 9, 8, 7, 6, 5, 3, 1], device='cuda:0')

8 9 8 7 6 5 3 1

step 561, loss 0.06572853773832321

Target: tensor([5, 0, 5, 0, 6, 0, 0, 2], device='cuda:0')

Pred: tensor([5, 0, 5, 0, 6, 0, 0, 2], device='cuda:0')

5 0 5 0 6 0 0 2

step 594, loss 0.00941784493625164

Target: tensor([1, 2, 6, 6, 0, 0, 3, 5], device='cuda:0')

Pred: tensor([1, 2, 6, 6, 0, 0, 3, 5], device='cuda:0')

1 2 6 6 0 0 3 5

step 627, loss 0.02076229266822338

Target: tensor([9, 1, 5, 6, 1, 9, 7, 2], device='cuda:0')

Pred: tensor([9, 1, 5, 6, 1, 9, 7, 2], device='cuda:0')

A horizontal strip showing the handwritten digits 9, 1, 5, 6, 1, 9, 7, and 2 in white on a black background.

step 660, loss 0.008764399215579033

Target: tensor([9, 8, 5, 5, 5, 8, 4, 3], device='cuda:0')

Pred: tensor([9, 8, 5, 5, 5, 8, 4, 3], device='cuda:0')

A horizontal strip showing the handwritten digits 9, 8, 5, 5, 5, 8, 4, and 3 in white on a black background.

step 693, loss 0.5224204063415527

Target: tensor([3, 8, 3, 6, 6, 4, 2, 4], device='cuda:0')

Pred: tensor([3, 8, 3, 6, 6, 4, 2, 4], device='cuda:0')

A horizontal strip showing the handwritten digits 3, 8, 3, 6, 6, 4, 2, and 4 in white on a black background.

step 726, loss 0.006300322245806456

Target: tensor([1, 6, 9, 7, 7, 1, 7, 9], device='cuda:0')

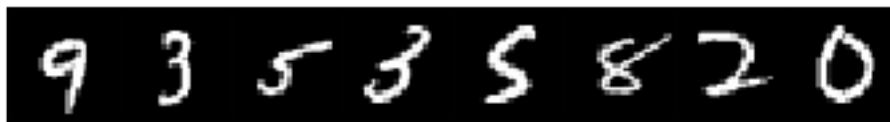
Pred: tensor([1, 6, 9, 7, 7, 1, 7, 9], device='cuda:0')

A horizontal strip showing the handwritten digits 1, 6, 9, 7, 7, 1, 7, and 9 in white on a black background.

step 759, loss 0.045768819749355316

Target: tensor([9, 3, 5, 3, 5, 8, 2, 0], device='cuda:0')

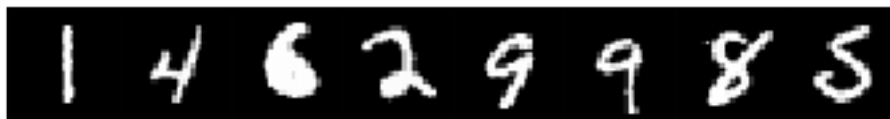
Pred: tensor([9, 3, 5, 3, 5, 8, 2, 0], device='cuda:0')

A horizontal strip showing the handwritten digits 9, 3, 5, 3, 5, 8, 2, and 0 in white on a black background.

step 792, loss 0.01284484937787056

Target: tensor([1, 4, 6, 2, 9, 9, 8, 5], device='cuda:0')

Pred: tensor([1, 4, 6, 2, 9, 9, 8, 5], device='cuda:0')

A horizontal strip showing the handwritten digits 1, 4, 6, 2, 9, 9, 8, and 5 in white on a black background.

step 825, loss 0.027232781052589417

Target: tensor([6, 4, 9, 3, 7, 2, 2, 2], device='cuda:0')

Pred: tensor([6, 4, 9, 3, 7, 2, 2, 2], device='cuda:0')

A horizontal strip showing the handwritten digits 6, 4, 9, 3, 7, 2, 2, and 2 in white on a black background.

step 858, loss 0.02522405982017517

Target: tensor([2, 4, 3, 7, 7, 8, 0, 8], device='cuda:0')

Pred: tensor([2, 4, 3, 7, 7, 8, 0, 8], device='cuda:0')

24377808

step 891, loss 0.016603466123342514

Target: tensor([1, 1, 4, 4, 0, 0, 5, 4], device='cuda:0')

Pred: tensor([1, 1, 4, 4, 0, 0, 5, 4], device='cuda:0')

11440054

step 924, loss 0.013803314417600632

Target: tensor([9, 7, 4, 4, 4, 1, 3, 7], device='cuda:0')

Pred: tensor([9, 7, 4, 4, 4, 1, 3, 7], device='cuda:0')

97444137

step 957, loss 0.005886902566999197

Target: tensor([7, 1, 0, 9, 0, 5, 2, 7], device='cuda:0')

Pred: tensor([7, 1, 0, 9, 0, 5, 2, 7], device='cuda:0')

71090527

step 990, loss 0.009861582890152931

Target: tensor([0, 9, 4, 2, 3, 2, 8, 1], device='cuda:0')

Pred: tensor([0, 9, 4, 2, 3, 2, 8, 1], device='cuda:0')



step 1023, loss 0.011333133094012737

Target: tensor([4, 3, 8, 7, 3, 9, 0, 8], device='cuda:0')

Pred: tensor([4, 3, 8, 7, 3, 9, 0, 8], device='cuda:0')



step 1056, loss 0.006470714695751667

Target: tensor([4, 1, 6, 3, 0, 6, 9, 0], device='cuda:0')

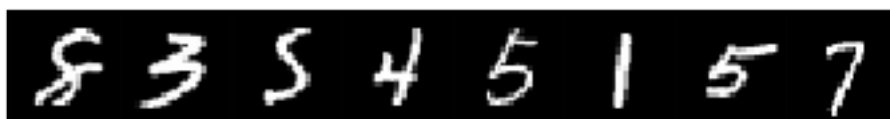
Pred: tensor([4, 1, 6, 3, 0, 6, 9, 0], device='cuda:0')



step 1089, loss 0.01293917279690504

Target: tensor([8, 3, 5, 4, 5, 1, 5, 7], device='cuda:0')

Pred: tensor([8, 3, 5, 4, 5, 1, 5, 7], device='cuda:0')



step 1122, loss 0.004658799152821302

Target: tensor([1, 1, 5, 7, 0, 3, 6, 2], device='cuda:0')

Pred: tensor([1, 1, 5, 7, 0, 3, 6, 2], device='cuda:0')

A horizontal strip showing the handwritten digits '1 1 5 7 0 3 6 2' in white on a black background.

step 1155, loss 0.005574997514486313

Target: tensor([4, 4, 9, 5, 3, 0, 5, 1], device='cuda:0')

Pred: tensor([4, 4, 9, 5, 3, 0, 5, 1], device='cuda:0')

A horizontal strip showing the handwritten digits '4 4 9 5 3 0 5 1' in white on a black background.

step 1188, loss 0.005826111882925034

Target: tensor([8, 9, 1, 6, 6, 4, 3, 3], device='cuda:0')

Pred: tensor([8, 9, 1, 6, 6, 4, 3, 3], device='cuda:0')

A horizontal strip showing the handwritten digits '8 9 1 6 6 4 3 3' in white on a black background.

step 1221, loss 0.007030262611806393

Target: tensor([9, 5, 7, 2, 9, 7, 7, 7], device='cuda:0')

Pred: tensor([9, 5, 7, 2, 9, 7, 7, 7], device='cuda:0')

9 5 7 2 9 7 7 7

step 1254, loss 0.008547605946660042

Target: tensor([1, 7, 2, 6, 8, 1, 0, 5], device='cuda:0')

Pred: tensor([1, 7, 2, 6, 8, 1, 0, 5], device='cuda:0')

1 7 2 6 8 1 0 5

step 1287, loss 0.002496434608474374

Target: tensor([8, 6, 1, 1, 5, 8, 7, 6], device='cuda:0')

Pred: tensor([8, 6, 1, 1, 5, 8, 7, 6], device='cuda:0')

8 6 1 1 5 8 7 6

step 1320, loss 0.003319466020911932

Target: tensor([0, 0, 6, 4, 8, 9, 8, 1], device='cuda:0')

Pred: tensor([0, 0, 6, 4, 8, 9, 8, 1], device='cuda:0')

0 0 6 4 8 9 8 1

step 1353, loss 0.0028432593680918217
Target: tensor([9, 2, 0, 5, 8, 7, 1, 4], device='cuda:0')
Pred: tensor([9, 2, 0, 5, 8, 7, 1, 4], device='cuda:0')

A black rectangular box containing the handwritten digits '9 2 0 5 8 7 1 4' in white ink.

step 1386, loss 0.003326366189867258
Target: tensor([9, 9, 3, 2, 0, 5, 9, 0], device='cuda:0')
Pred: tensor([9, 9, 3, 2, 0, 5, 9, 0], device='cuda:0')

A black rectangular box containing the handwritten digits '9 9 3 2 0 5 9 0' in white ink.

step 1419, loss 0.00521630747243762
Target: tensor([5, 8, 5, 7, 1, 1, 6, 5], device='cuda:0')
Pred: tensor([5, 8, 5, 7, 1, 1, 6, 5], device='cuda:0')

A black rectangular box containing the handwritten digits '5 8 5 7 1 1 6 5' in white ink.

step 1452, loss 0.003354680724442005
Target: tensor([5, 3, 7, 9, 4, 8, 4, 6], device='cuda:0')
Pred: tensor([5, 3, 7, 9, 4, 8, 4, 6], device='cuda:0')

A black rectangular box containing the handwritten digits '5 3 7 9 4 8 4 6' in white ink.

step 1485, loss 0.009109675884246826

Target: tensor([6, 8, 5, 2, 9, 3, 4, 5], device='cuda:0')

Pred: tensor([6, 8, 5, 2, 9, 3, 4, 5], device='cuda:0')

A black rectangular box containing the handwritten digits '6 8 5 2 9 3 4 5' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 1518, loss 0.006141078658401966

Target: tensor([0, 7, 2, 6, 1, 9, 7, 7], device='cuda:0')

Pred: tensor([0, 7, 2, 6, 1, 9, 7, 7], device='cuda:0')

A black rectangular box containing the handwritten digits '0 7 2 6 1 9 7 7' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 1551, loss 0.0018050407525151968

Target: tensor([6, 2, 7, 1, 1, 8, 3, 7], device='cuda:0')

Pred: tensor([6, 2, 7, 1, 1, 8, 3, 7], device='cuda:0')

A black rectangular box containing the handwritten digits '6 2 7 1 1 8 3 7' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 1584, loss 0.00226784311234951

Target: tensor([7, 7, 4, 2, 3, 8, 2, 8], device='cuda:0')

Pred: tensor([7, 7, 4, 2, 3, 8, 2, 8], device='cuda:0')

1 7 4 2 3 8 2 8

step 1617, loss 0.0015142107149586082

Target: tensor([3, 3, 1, 8, 5, 7, 7, 3], device='cuda:0')

Pred: tensor([3, 3, 1, 8, 5, 7, 7, 3], device='cuda:0')

3 3 1 8 5 7 7 3

step 1650, loss 0.005291391164064407

Target: tensor([6, 1, 5, 4, 4, 7, 2, 9], device='cuda:0')

Pred: tensor([6, 1, 5, 4, 4, 7, 2, 9], device='cuda:0')

6 1 5 4 4 7 2 9

step 1683, loss 0.0010156281059607863

Target: tensor([1, 1, 3, 6, 6, 5, 2, 6], device='cuda:0')

Pred: tensor([1, 1, 3, 6, 6, 5, 2, 6], device='cuda:0')

1 1 3 6 6 5 2 6

step 1716, loss 0.0018912971718236804

Target: tensor([0, 5, 0, 6, 3, 1, 9, 7], device='cuda:0')

Pred: tensor([0, 5, 0, 6, 3, 1, 9, 7], device='cuda:0')

A black rectangular box containing the handwritten digits '0 5 0 6 3 1 9 7' in white ink.

step 1749, loss 0.003080256748944521

Target: tensor([3, 9, 5, 9, 4, 1, 9, 2], device='cuda:0')

Pred: tensor([3, 9, 5, 9, 4, 1, 9, 2], device='cuda:0')

A black rectangular box containing the handwritten digits '3 9 5 9 4 1 9 2' in white ink.

step 1782, loss 0.0018953669350594282

Target: tensor([1, 3, 4, 8, 7, 5, 6, 0], device='cuda:0')

Pred: tensor([1, 3, 4, 8, 7, 5, 6, 0], device='cuda:0')

A black rectangular box containing the handwritten digits '1 3 4 8 7 5 6 0' in white ink.

step 1815, loss 0.0012907194904983044

Target: tensor([3, 0, 3, 2, 8, 0, 0, 1], device='cuda:0')

Pred: tensor([3, 0, 3, 2, 8, 0, 0, 1], device='cuda:0')

A black rectangular box containing the handwritten digits '3 0 3 2 8 0 0 1' in white ink.

step 1848, loss 0.004156704992055893

Target: tensor([0, 2, 7, 0, 6, 2, 0, 4], device='cuda:0')

Pred: tensor([0, 2, 7, 0, 6, 2, 0, 4], device='cuda:0')

A horizontal strip showing the handwritten digits '0 2 7 0 6 2 0 4' in white on a black background.

step 1881, loss 0.0035496680065989494

Target: tensor([2, 3, 4, 7, 4, 8, 6, 4], device='cuda:0')

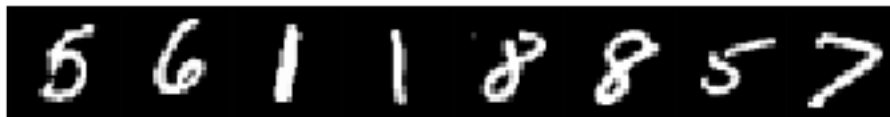
Pred: tensor([2, 3, 4, 7, 4, 8, 6, 4], device='cuda:0')

A horizontal strip showing the handwritten digits '2 3 4 7 4 8 6 4' in white on a black background.

step 1914, loss 0.0014311763225123286

Target: tensor([5, 6, 1, 1, 8, 8, 5, 7], device='cuda:0')

Pred: tensor([5, 6, 1, 1, 8, 8, 5, 7], device='cuda:0')

A horizontal strip showing the handwritten digits '5 6 1 1 8 8 5 7' in white on a black background.

step 1947, loss 0.0030275534372776747

Target: tensor([7, 3, 6, 7, 8, 0, 3, 3], device='cuda:0')

Pred: tensor([7, 3, 6, 7, 8, 0, 3, 3], device='cuda:0')

7 3 6 7 8 0 3 3

step 1980, loss 0.0022723025176674128

Target: tensor([9, 6, 7, 4, 6, 1, 5, 8], device='cuda:0')

Pred: tensor([9, 6, 7, 4, 6, 1, 5, 8], device='cuda:0')

9 6 7 4 6 1 5 8

step 2013, loss 0.004244609735906124

Target: tensor([4, 9, 5, 1, 8, 8, 5, 8], device='cuda:0')

Pred: tensor([4, 9, 5, 1, 8, 8, 5, 8], device='cuda:0')

4 9 5 1 8 8 5 8

step 2046, loss 0.0032289219088852406

Target: tensor([3, 5, 9, 8, 8, 7, 3, 7], device='cuda:0')

Pred: tensor([3, 5, 9, 8, 8, 7, 3, 7], device='cuda:0')

3 5 9 8 8 7 3 7

step 2079, loss 0.0012416639365255833

Target: tensor([6, 1, 4, 7, 0, 3, 2, 6], device='cuda:0')

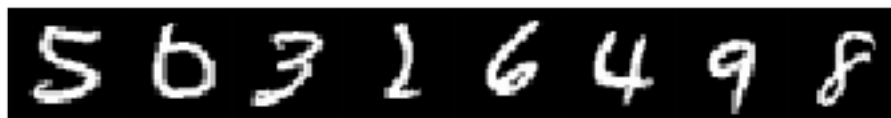
Pred: tensor([6, 1, 4, 7, 0, 3, 2, 6], device='cuda:0')

A horizontal strip showing the handwritten digits 6, 1, 4, 7, 0, 3, 2, and 6 in white on a black background.

step 2112, loss 0.0020870317239314318

Target: tensor([5, 6, 3, 2, 6, 4, 9, 8], device='cuda:0')

Pred: tensor([5, 6, 3, 2, 6, 4, 9, 8], device='cuda:0')

A horizontal strip showing the handwritten digits 5, 6, 3, 2, 6, 4, 9, and 8 in white on a black background.

step 2145, loss 0.002490028738975525

Target: tensor([2, 9, 8, 2, 3, 3, 5, 5], device='cuda:0')

Pred: tensor([2, 9, 8, 2, 3, 3, 5, 5], device='cuda:0')

A horizontal strip showing the handwritten digits 2, 9, 8, 2, 3, 3, 5, and 5 in white on a black background.

step 2178, loss 0.0013264240697026253

Target: tensor([9, 5, 2, 4, 6, 1, 8, 4], device='cuda:0')

Pred: tensor([9, 5, 2, 4, 6, 1, 8, 4], device='cuda:0')

A horizontal strip showing the handwritten digits 9, 5, 2, 4, 6, 1, 8, and 4 in white on a black background.

step 2211, loss 0.0009308563312515616

Target: tensor([6, 1, 4, 7, 4, 3, 3, 2], device='cuda:0')

Pred: tensor([6, 1, 4, 7, 4, 3, 3, 2], device='cuda:0')

A horizontal strip showing the handwritten digits 6, 1, 4, 7, 4, 3, 3, 2 in white on a black background.

step 2244, loss 0.0012827360769733787

Target: tensor([9, 4, 1, 9, 0, 0, 9, 2], device='cuda:0')

Pred: tensor([9, 4, 1, 9, 0, 0, 9, 2], device='cuda:0')

A horizontal strip showing the handwritten digits 9, 4, 1, 9, 0, 0, 9, 2 in white on a black background.

step 2277, loss 0.001407506293617189

Target: tensor([7, 3, 7, 3, 3, 5, 1, 1], device='cuda:0')

Pred: tensor([7, 3, 7, 3, 3, 5, 1, 1], device='cuda:0')

A horizontal strip showing the handwritten digits 7, 3, 7, 3, 3, 5, 1, 1 in white on a black background.

step 2310, loss 0.002045449335128069

Target: tensor([4, 7, 7, 1, 3, 9, 0, 6], device='cuda:0')

Pred: tensor([4, 7, 7, 1, 3, 9, 0, 6], device='cuda:0')

4 7 7 1 3 9 0 6

step 2343, loss 0.002122690202668309

Target: tensor([6, 8, 7, 6, 6, 2, 4, 8], device='cuda:0')

Pred: tensor([6, 8, 7, 6, 6, 2, 4, 8], device='cuda:0')

6 8 7 6 6 2 4 8

step 2376, loss 0.0005678567686118186

Target: tensor([5, 4, 8, 1, 6, 3, 4, 6], device='cuda:0')

Pred: tensor([5, 4, 8, 1, 6, 3, 4, 6], device='cuda:0')

5 4 8 1 6 3 4 6

step 2409, loss 0.0014561372809112072

Target: tensor([9, 9, 6, 6, 8, 4, 8, 1], device='cuda:0')

Pred: tensor([9, 9, 6, 6, 8, 4, 8, 1], device='cuda:0')

9 9 6 6 8 4 8 1

step 2442, loss 0.0024195618461817503

Target: tensor([3, 3, 2, 7, 7, 2, 4, 7], device='cuda:0')

Pred: tensor([3, 3, 2, 7, 7, 2, 4, 7], device='cuda:0')

A horizontal strip of a black background with white handwritten digits. The digits are '3', '3', '2', '7', '7', '2', '4', and '7' from left to right.

step 2475, loss 0.0008656043792143464

Target: tensor([1, 8, 1, 4, 5, 5, 6, 3], device='cuda:0')

Pred: tensor([1, 8, 1, 4, 5, 5, 6, 3], device='cuda:0')

A horizontal strip of a black background with white handwritten digits. The digits are '1', '8', '1', '4', '5', '5', '6', and '3' from left to right.

step 2508, loss 0.00240669259801507

Target: tensor([8, 8, 1, 2, 6, 9, 2, 6], device='cuda:0')

Pred: tensor([8, 8, 1, 2, 6, 9, 2, 6], device='cuda:0')

A horizontal strip of a black background with white handwritten digits. The digits are '8', '8', '1', '2', '6', '9', '2', and '6' from left to right.

step 2541, loss 0.002071989234536886

Target: tensor([8, 1, 2, 7, 4, 7, 0, 2], device='cuda:0')

Pred: tensor([8, 1, 2, 7, 4, 7, 0, 2], device='cuda:0')

A horizontal strip of a black background with white handwritten digits. The digits are '8', '1', '2', '7', '4', '7', '0', and '2' from left to right.

step 2574, loss 0.0026595816016197205

Target: tensor([6, 8, 1, 0, 9, 6, 3, 0], device='cuda:0')

Pred: tensor([6, 8, 1, 0, 9, 6, 3, 0], device='cuda:0')

A horizontal strip showing the handwritten digits '6 8 1 0 9 6 3 0' in white on a black background.

step 2607, loss 0.001836067414842546

Target: tensor([4, 1, 3, 8, 7, 3, 5, 2], device='cuda:0')

Pred: tensor([4, 1, 3, 8, 7, 3, 5, 2], device='cuda:0')

A horizontal strip showing the handwritten digits '4 1 3 8 7 3 5 2' in white on a black background.

step 2640, loss 0.0023896931670606136

Target: tensor([6, 3, 5, 0, 8, 9, 0, 9], device='cuda:0')

Pred: tensor([6, 3, 5, 0, 8, 9, 0, 9], device='cuda:0')

A horizontal strip showing the handwritten digits '6 3 5 0 8 9 0 9' in white on a black background.

step 2673, loss 0.0008393566822633147

Target: tensor([4, 2, 7, 8, 1, 5, 5, 7], device='cuda:0')

Pred: tensor([4, 2, 7, 8, 1, 5, 5, 7], device='cuda:0')

4 2 7 8 1 5 5 7

step 2706, loss 0.0008002709364518523

Target: tensor([9, 9, 6, 4, 4, 9, 5, 5], device='cuda:0')

Pred: tensor([9, 9, 6, 4, 4, 9, 5, 5], device='cuda:0')

9 9 6 4 4 9 5 5

step 2739, loss 0.0009052536333911121

Target: tensor([0, 3, 5, 9, 7, 0, 7, 1], device='cuda:0')

Pred: tensor([0, 3, 5, 9, 7, 0, 7, 1], device='cuda:0')

0 3 5 9 7 0 7 1

step 2772, loss 0.00284338416531682

Target: tensor([0, 2, 4, 1, 7, 2, 8, 3], device='cuda:0')


Pred: tensor([0, 2, 4, 1, 7, 2, 8, 3], device='cuda:0')

0 2 4 1 7 2 8 3

step 2805, loss 0.001619724789634347

Target: tensor([1, 1, 9, 4, 6, 7, 3, 2], device='cuda:0')

Pred: tensor([1, 1, 9, 4, 6, 7, 3, 2], device='cuda:0')

A black rectangular box containing the handwritten digits '1 1 9 4 6 7 3 2' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 2838, loss 0.002499950584024191

Target: tensor([1, 5, 6, 9, 5, 8, 5, 3], device='cuda:0')

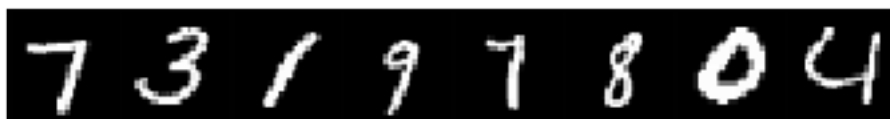
Pred: tensor([1, 5, 6, 9, 5, 8, 5, 3], device='cuda:0')

A black rectangular box containing the handwritten digits '1 5 6 9 5 8 5 3' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 2871, loss 0.0008655176497995853

Target: tensor([7, 3, 1, 9, 7, 8, 0, 4], device='cuda:0')

Pred: tensor([7, 3, 1, 9, 7, 8, 0, 4], device='cuda:0')

A black rectangular box containing the handwritten digits '7 3 1 9 7 8 0 4' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 2904, loss 0.001871989807114005

Target: tensor([7, 5, 5, 3, 4, 1, 0, 0], device='cuda:0')

Pred: tensor([7, 5, 5, 3, 4, 1, 0, 0], device='cuda:0')

A black rectangular box containing the handwritten digits '7 5 5 3 4 1 0 0' in white ink. The digits are slightly slanted and have a casual, handwritten style.

step 2937, loss 0.0018203234067186713

Target: tensor([5, 9, 1, 6, 6, 3, 7, 4], device='cuda:0')

Pred: tensor([5, 9, 1, 6, 6, 3, 7, 4], device='cuda:0')

A black rectangular box containing the handwritten digits '5 9 1 6 6 3 7 4' in white ink. The digits are slightly blurred and have a casual, handwritten style.

step 2970, loss 0.002758538816124201

Target: tensor([0, 8, 4, 2, 9, 1, 0, 3], device='cuda:0')

Pred: tensor([0, 8, 4, 2, 9, 1, 0, 3], device='cuda:0')

A black rectangular box containing the handwritten digits '0 8 4 2 9 1 0 3' in white ink. The digits are slightly blurred and have a casual, handwritten style.

step 3003, loss 0.0012137601152062416

Target: tensor([4, 9, 7, 7, 3, 8, 5, 9], device='cuda:0')

Pred: tensor([4, 9, 7, 7, 3, 8, 5, 9], device='cuda:0')

A black rectangular box containing the handwritten digits '4 9 7 7 3 8 5 9' in white ink. The digits are slightly blurred and have a casual, handwritten style.

step 3036, loss 0.00028281507547944784

Target: tensor([2, 1, 2, 4, 4, 4, 5, 2], device='cuda:0')

Pred: tensor([2, 1, 2, 4, 4, 4, 5, 2], device='cuda:0')

2 1 2 4 4 4 5 2

step 3069, loss 0.0020678387954831123

Target: tensor([7, 4, 7, 0, 1, 1, 0, 8], device='cuda:0')

Pred: tensor([7, 4, 7, 0, 1, 1, 0, 8], device='cuda:0')

7 4 7 0 1 1 0 8

step 3102, loss 0.0013928182888776064

Target: tensor([1, 8, 8, 6, 7, 9, 7, 9], device='cuda:0')

Pred: tensor([1, 8, 8, 6, 7, 9, 7, 9], device='cuda:0')

1 8 8 6 7 9 7 9

step 3135, loss 0.0006481488817371428

Target: tensor([8, 5, 0, 7, 3, 5, 6, 4], device='cuda:0')


Pred: tensor([8, 5, 0, 7, 3, 5, 6, 4], device='cuda:0')

8 5 0 7 3 5 6 4

step 3168, loss 0.001968739088624716

Target: tensor([2, 4, 7, 0, 2, 2, 4, 9], device='cuda:0')

Pred: tensor([2, 4, 7, 0, 2, 2, 4, 9], device='cuda:0')

A horizontal strip showing the handwritten digits '2 4 7 0 2 2 4 9' in white on a black background.

step 3201, loss 0.0011353774461895227

Target: tensor([8, 3, 9, 3, 4, 7, 3, 2], device='cuda:0')

Pred: tensor([8, 3, 9, 3, 4, 7, 3, 2], device='cuda:0')

A horizontal strip showing the handwritten digits '8 3 9 3 4 7 3 2' in white on a black background.

step 3234, loss 0.0007010146509855986

Target: tensor([4, 3, 5, 4, 0, 3, 7, 5], device='cuda:0')

Pred: tensor([4, 3, 5, 4, 0, 3, 7, 5], device='cuda:0')



A horizontal strip showing the handwritten digits '4 3 5 4 0 3 7 5' in white on a black background.

step 3267, loss 0.00040904409252107143

Target: tensor([3, 5, 1, 2, 8, 1, 8, 8], device='cuda:0')

Pred: tensor([3, 5, 1, 2, 8, 1, 8, 8], device='cuda:0')

A horizontal strip showing the handwritten digits '3 5 1 2 8 1 8 8' in white on a black background.



1.10 Guide: Model Saving and Loading

Once the training is completed, we can save the model on disk for evaluation or future use. It is also helpful to save the model regularly in case of unexpected situations. Below is the snippet of how to save and load a model. More information can be found [here](#). You will be asked to evaluate your LeNet in the end.

```
# Saving a model on disk:
torch.save(net.state_dict(), PATH_to_save)

# Loading a model from disk:
net = LeNet(img_c)
net.load_state_dict(torch.load(PATH_to_save))
```

1.11 Task 5: Weight Decay

As we can see, LeNet now is overfitted on this 1k dataset. Instead of providing more data, we can use Weight Decay to improve its generalization ability.

To-dos: - (5 points) Write the dataset loading script, training script - (5 points) Add **weights penalty** into the loss function. - (10 points) Train LeNet from scratch. - (10 points) Plot out the training loss curve, validation loss curve, training accuracy curve, and validation accuracy curve as those plots in the above section.

```
[23]: def weight_decay(loss, weights, lamb, batch):
        # add the weight decay penalty to the loss
        penalty = torch.tensor(0.0).cuda()
        num_w = torch.tensor(0.0).cuda()
        i=0
        for w in weights:
            if i%2==0:
                penalty += (w*w).sum()
                num_w += w.clone().reshape(1,-1).shape[1]
            i+=1
        l = torch.tensor(0.5*lamb).cuda()
        loss += l*penalty/torch.tensor(batch).float().cuda()
        return loss
```

```
[18]: #####
        # Dataset loading script
        dataset_size = 1000
        validation_size = int(0.2 * 1000)

        ds = MNISTDataset(ds_size=dataset_size)

        # split the dataset into training set and validation set
        train_ds, val_ds = torch.utils.data.random_split(ds, [dataset_size -
        ↪ validation_size, validation_size])
```

```

# training batch size, hyper-parameter
batch_size = 24

# dataset loader
train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
    ↪drop_last=True)
val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=True, drop_last=True)
img_channel = 1

#####
# Training script
# learning rate, hyper-parameter
lr = 1e-3

# using GPU if it's available
device = torch.device('cuda')

net = LeNet(img_channel, device)
# keep a list of model parameters
params = [p for (n, p) in net.named_parameters()]

# training epochs, hyper-parameter
epochs = 100

# keep tracking of the changing of loss and accuracy of predictions
train_loss_list = []
val_loss_list = []
train_acc_list = []
val_acc_list = []
print_interval = 350
for e in range(epochs):
    net.train()
    for i, (imgs, lbs) in enumerate(train_dl):
        imgs = imgs.to(device)
        lbs = lbs.to(device)
        loss, prob = net(imgs, lbs)
        # add the weight decay penalty to the loss
        lamb = 2.0
        loss = weight_decay(loss, params, lamb, lbs.shape[0])

    net.zero_grad()
    grad = torch.autograd.grad(loss, params)

    # update weights
    step(params, grad, lr)

    # obtain the predictions

```

```

pred = prob.argmax(dim=-1).view(batch_size)
acc = (pred == lbs).float().mean()

if (i + e*len(train_dl))% print_interval == 0:
    print("e {}, step {}, loss {}".format( e, i + e*len(train_dl), loss))
    print("Target:\t {} \nPred:\t {}".format(lbs[:8], pred[:8]))
    # visualize some samples
    imgs_to_vis = vutils.make_grid(imgs[:8].cpu()+125., nrow=8, pad_value=1)
    plt.imshow(imgs_to_vis.permute(1,2,0).numpy().astype(np.uint8))
    plt.axis("off")
    plt.show()

train_loss_list.append(loss.detach().mean())
train_acc_list.append(acc.detach().mean())

net.eval()
for i, (imgs, lbs) in enumerate(val_dl):
    imgs = imgs.to(device)
    lbs = lbs.to(device)
    loss, prob = net(imgs, lbs)
    loss = weight_decay(loss, params, lamb, lbs.shape[0])

    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()

val_loss_list.append(loss.detach().mean())
val_acc_list.append(acc.detach().mean())

# plotting logs
plt.plot(np.arange(epochs), train_loss_list, '-r',
         np.arange(epochs), val_loss_list, '-g')
plt.legend(('training error', 'validation error'))
plt.show()
plt.plot(np.arange(epochs), train_acc_list, '-r',
         np.arange(epochs), val_acc_list, '-g')
plt.legend(('training acc', 'validation acc'))
plt.show()

```

```


e 0, step 0, loss 44.088890075683594
Target:  tensor([8, 6, 7, 6, 6, 9, 9, 1], device='cuda:0')
Pred:    tensor([7, 7, 2, 7, 7, 2, 7, 7], device='cuda:0')


```

8 6 7 6 6 9 9 1

e 60, step 2000, loss 18.280153274536133
Target: tensor([9, 6, 1, 8, 7, 1, 1, 5], device='cuda:0')
Pred: tensor([9, 6, 1, 8, 7, 1, 1, 5], device='cuda:0')

9 6 1 8 7 1 1 5





1.12 Task 6: Dropout

Dropout can also help with generalization. It randomly drops units from the neural network during training. Thus, training a neural network with dropout can be seen as training a collection of many sub-networks.

To-dos: - (5 points) Complete the `dropout_forward()` function for class `Dropout`. - (5 points) Modify the class `LeNet` with dropout layer and build a new class named `LeNetDrop`: add a dropout layer after layer `c1`, layer `c3`, and layer `c5`. - Write the dataset loading script, training script. - Train `LeNetDrop` on the 1k dataset. - (10 points) Plot out the training loss curve, validation loss curve, training accuracy curve and validation accuracy curve.

```
[17]: class Dropout(nn.Module):
    def __init__(self, p, device):
        super(Dropout, self).__init__()
        """
        inputs:
            p: scalar, the probability of an element to be zeroed.
        """
        self.p = p
        self.device = device

    def dropout_forward(self, X, training=True):
        """
```

```

    inputs:
        X: (N, *), input of dropout layer
        training: boolean, apply dropout if true. Note: We do not apply dropout,
        ↪ during testing.
    outputs:
        Y: (N, *)
    """
    ### Start the code here ###
    if training == True:
        dr = (torch.rand(X.shape) > self.p).cuda().float()
        Y = X*dr
        Y = Y/(1.0-self.p)
    else:
        Y = X

    ### End of the code ###
    return Y

def forward(self, X, training=True):
    return self.dropout_forward(X, training)

```

```

[18]: # dataset loading, training script and other codes starts here:
# LeNetDrop sketch code
class LeNetDrop(nn.Module):
    def __init__(self, img_c, device):
        super(LeNetDrop, self).__init__()
        drop_p = 0.2
        self.c1 = Conv2D(img_c, 6, [5,5], [1,1], [2,2,2,2], device)
        self.d1 = Dropout(drop_p, device)
        self.p2 = nn.MaxPool2d(2, stride=2)
        self.c3 = Conv2D(6, 16, [5,5], [1,1], [0,0,0,0], device)
        self.d3 = Dropout(drop_p, device)
        self.p4 = nn.MaxPool2d(2, stride=2)
        self.f5 = Linear(400, 120, device)
        self.d5 = Dropout(drop_p, device)
        self.f6 = Linear(120, 84, device)
        self.f7 = Linear(84, 10, device)
        self.device = device

    def forward(self, imgs, labels, training=True):
        """
        inputs:
            imgs: (N, C, H, W), training samples from the MNIST training set, where N
            ↪ is the number of samples (batch_size),
            C is the image color channle number, H and W are the spatial size of
            ↪ the input images.

```

```

        labels: (N, L), ground truth for the input images, where N is the number
        ↪ of samples (batch_size) and L is the
            number of classes.
    outputs:
        loss: (1,), mean loss value over this batch of inputs.

    """
    N = imgs.shape[0]

    o_c1 = F.relu(self.c1(imgs))
    o_d1 = self.d1(o_c1, training)
    o_p2 = self.p2(o_d1)
    o_c3 = F.relu(self.c3(o_p2))
    o_d3 = self.d3(o_c3, training)
    o_p4 = self.p4(o_d3)

    ### Start the code here ###
    # 1. Please complete the rest of LeNet to get the scores predicted by LeNet
    ↪ for each input images #
    # need to flatten the matrix before forwarding to the dense layer
    o_f5 = F.relu(self.f5(o_p4.reshape(o_p4.shape[0], -1)))
    o_d5 = self.d5(o_f5, training)
    o_f6 = F.relu(self.f6(o_d5))
    o_f7 = self.f7(o_f6)
    # 2. Please use the implemented objective function to obtain the losses of
    ↪ each input. #
    p = softmax1d(o_f7)
    # 3. We will return the mean value of the losses. #
    loss = cross_entropy_loss(p, labels)
    ### End of the code ###
    return loss.mean(), p

```

[21]:

```

#####
# Train LeNetDrop on the 1k dataset
#####
# Dataset loading script
dataset_size = 1000
validation_size = int(0.2 * 1000)

ds = MNISTDataset(ds_size=dataset_size)

# split the dataset into training set and validation set
train_ds, val_ds = torch.utils.data.random_split(ds, [dataset_size -
    ↪ validation_size, validation_size])

# training batch size, hyper-parameter
batch_size = 24

```

```

# dataset loader
train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
    ↳drop_last=True)
val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=True, drop_last=True)
img_channel = 1

#####
# Training script
# learning rate, hyper-parameter
lr = 1e-3

# using GPU if it's available
device = torch.device('cuda')

net = LeNetDrop(img_channel, device)
# keep a list of model parameters
params = [p for (n, p) in net.named_parameters()]

# training epochs, hyper-parameter
epochs = 100

# keep tracking of the changing of loss and accuracy of predictions
train_loss_list = []
val_loss_list = []
train_acc_list = []
val_acc_list = []
print_interval = 350
for e in range(epochs):
    net.train()
    for i, (imgs, lbs) in enumerate(train_dl):
        imgs = imgs.to(device)
        lbs = lbs.to(device)
        loss, prob = net(imgs, lbs, training=True)
        # add the weight decay penalty to the loss
        lamb = 2.0
        loss = weight_decay(loss, params, lamb, lbs.shape[0])

    net.zero_grad()
    grad = torch.autograd.grad(loss, params)

    # update weights
    step(params, grad, lr)

    # obtain the predictions
    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()

```



```

if (i + e*len(train_dl)) % print_interval == 0:
    print("e {}, step {}, loss {}".format( e, i + e*len(train_dl), loss))
    print("Target:\t {} \nPred:\t {}".format(lbs[:8], pred[:8]))
    # visualize some samples
    imgs_to_vis = vutils.make_grid(imgs[:8].cpu()+125., nrow=8, pad_value=1)
    plt.imshow(imgs_to_vis.permute(1,2,0).numpy().astype(np.uint8))
    plt.axis("off")
    plt.show()
train_loss_list.append(loss.detach().mean())
train_acc_list.append(acc.detach().mean())

net.eval()
for i, (imgs, lbs) in enumerate(val_dl):
    imgs = imgs.to(device)
    lbs = lbs.to(device)
    loss, prob = net(imgs, lbs, False)
    # add the weight decay penalty to the loss
    lamb = 2.0
    loss = weight_decay(loss, params, lamb, lbs.shape[0])

    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()

val_loss_list.append(loss.detach().mean())
val_acc_list.append(acc.detach().mean())

# plotting logs
plt.plot(np.arange(epochs), train_loss_list, '-r',
         np.arange(epochs), val_loss_list, '-g')
plt.legend(('training error', 'validation error'))
plt.show()
plt.plot(np.arange(epochs), train_acc_list, '-r',
         np.arange(epochs), val_acc_list, '-g')
plt.legend(('training acc', 'validation acc'))
plt.show()

```

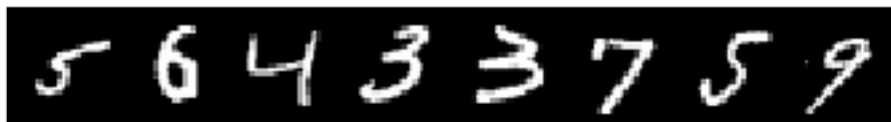
e 0, step 0, loss 40.31529235839844

Target: tensor([6, 3, 4, 0, 9, 9, 4, 9], device='cuda:0')

Pred: tensor([6, 0, 7, 0, 6, 9, 0, 6], device='cuda:0')



e 10, step 350, loss 31.059011459350586
Target: tensor([5, 6, 4, 3, 3, 7, 5, 9], device='cuda:0')
Pred: tensor([5, 6, 4, 3, 3, 7, 5, 7], device='cuda:0')

A black rectangular box containing the handwritten digits '5 6 4 3 3 7 5 9' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 21, step 700, loss 23.398168563842773
Target: tensor([5, 0, 7, 1, 4, 3, 5, 2], device='cuda:0')
Pred: tensor([5, 0, 7, 1, 4, 3, 5, 2], device='cuda:0')

A black rectangular box containing the handwritten digits '5 0 7 1 4 3 5 2' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 31, step 1050, loss 22.58089256286621
Target: tensor([2, 1, 7, 2, 9, 3, 1, 5], device='cuda:0')
Pred: tensor([2, 1, 7, 0, 9, 0, 5, 5], device='cuda:0')

A black rectangular box containing the handwritten digits '2 1 7 0 9 3 1 5' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 42, step 1400, loss 20.371061325073242
Target: tensor([3, 0, 8, 8, 4, 8, 5, 9], device='cuda:0')
Pred: tensor([3, 0, 8, 8, 4, 6, 5, 9], device='cuda:0')

3 0 8 8 4 8 5 9

e 53, step 1750, loss 19.366426467895508

Target: tensor([9, 4, 7, 6, 4, 6, 3, 5], device='cuda:0')

Pred: tensor([7, 4, 7, 6, 4, 6, 3, 5], device='cuda:0')

9 4 7 6 4 6 3 5

e 63, step 2100, loss 18.303951263427734

Target: tensor([7, 2, 7, 4, 2, 1, 9, 0], device='cuda:0')

Pred: tensor([7, 2, 7, 4, 2, 1, 9, 0], device='cuda:0')

7 2 7 4 2 1 9 0

e 74, step 2450, loss 16.978912353515625

Target: tensor([7, 7, 0, 7, 8, 9, 1, 4], device='cuda:0')

Pred: tensor([7, 7, 0, 7, 8, 9, 1, 4], device='cuda:0')

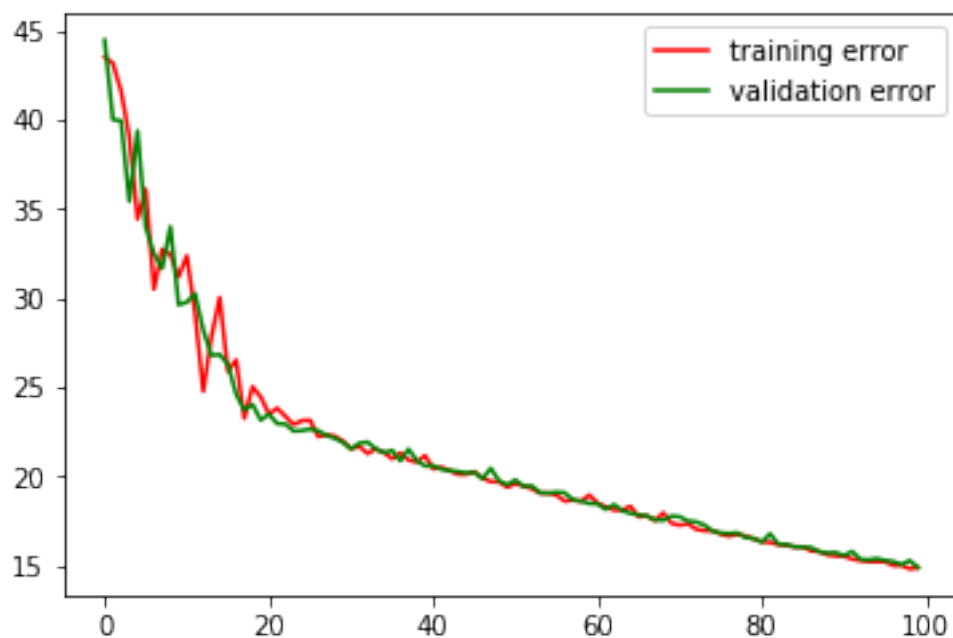
7 7 0 1 8 9 1 4


```
e 84, step 2800, loss 16.1416072845459
Target:  tensor([8, 8, 4, 2, 8, 5, 6, 2], device='cuda:0')
Pred:    tensor([8, 5, 4, 2, 8, 5, 6, 2], device='cuda:0')
```

8 8 4 2 8 5 6 2

```
e 95, step 3150, loss 15.568328857421875
Target:  tensor([5, 1, 6, 0, 2, 9, 7, 2], device='cuda:0')
Pred:    tensor([5, 1, 6, 0, 2, 9, 7, 2], device='cuda:0')
```

5 1 6 0 2 9 7 2





1.13 Task 7: Train on full dataset

Now, train the LeNetDrop on the full MNIST dataset.

To-do: - Write the dataset loading script, training script for LeNetDrop. - Train LeNetDrop on the full MNIST dataset below. - (10 points) Plot out the training loss curve, validation loss curve, training accuracy curve and validation accuracy curve.

```
[19]: #####
# Train LeNetDrop on the full dataset
#####
# Dataset loading script
train_ds = MNISTDataset(60000, split = "training")
val_ds = MNISTDataset(10000, split = "validation")

# training batch size, hyper-parameter
batch_size = 24

# dataset loader
train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
    ↳drop_last=True)
val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=True, drop_last=True)
img_channel = 1
```

[24]:

```
#####  
# Training script  
# learning rate, hyper-parameter  
lr = 1e-3  
  
# using GPU if it's available  
device = torch.device('cuda')  
  
net = LeNetDrop(img_channel, device)  
# keep a list of model parameters  
params = [p for (n, p) in net.named_parameters()]  
  
# training epochs, hyper-parameter  
epochs = 100  
  
# keep tracking of the changing of loss and accuracy of predictions  
train_loss_list = []  
val_loss_list = []  
train_acc_list = []  
val_acc_list = []  
print_interval = 2000  
for e in range(epochs):  
    net.train()  
    for i, (imgs, lbs) in enumerate(train_dl):  
        imgs = imgs.to(device)  
        lbs = lbs.to(device)  
        loss, prob = net(imgs, lbs, training=True)  
        # add the weight decay penalty to the loss  
        lamb = 2.0  
        loss = weight_decay(loss, params, lamb, lbs.shape[0])  
  
        net.zero_grad()  
        grad = torch.autograd.grad(loss, params)  
  
        # update weights  
        step(params, grad, lr)  
  
        # obtain the predictions  
        pred = prob.argmax(dim=-1).view(batch_size)  
        acc = (pred == lbs).float().mean()  
        if (i + e*len(train_dl)) % print_interval == 0:  
            print("e {}, step {}, loss {}".format(e, i + e*len(train_dl), loss))  
            print("Target:\t {} \n Pred:\t {}".format(lbs[:8], pred[:8]))  
            # visualize some samples  
            imgs_to_vis = vutils.make_grid(imgs[:8].cpu()+125., nrow=8, pad_value=1)  
            plt.imshow(imgs_to_vis.permute(1,2,0).numpy().astype(np.uint8))  
            plt.axis("off")
```

```

        plt.show()
    train_loss_list.append(loss.detach().mean())
    train_acc_list.append(acc.detach().mean())

    net.eval()
    for i, (imgs, lbs) in enumerate(val_dl):
        imgs = imgs.to(device)
        lbs = lbs.to(device)
        loss, prob = net(imgs, lbs, False)
        # add the weight decay penalty to the loss
        lamb = 2.0
        loss = weight_decay(loss, params, lamb, lbs.shape[0])

        pred = prob.argmax(dim=-1).view(batch_size)
        acc = (pred == lbs).float().mean()

    val_loss_list.append(loss.detach().mean())
    val_acc_list.append(acc.detach().mean())

# plotting logs
    plt.plot(np.arange(epochs), train_loss_list, '-r',
             np.arange(epochs), val_loss_list, '-g')
    plt.legend(('training error', 'validation error'))
    plt.show()
    plt.plot(np.arange(epochs), train_acc_list, '-r',
             np.arange(epochs), val_acc_list, '-g')
    plt.legend(('training acc', 'validation acc'))
    plt.show()

```

e 0, step 0, loss 45.258338928222656

Target: tensor([3, 1, 4, 7, 1, 9, 5, 4], device='cuda:0')

Pred: tensor([0, 8, 0, 4, 0, 4, 8, 0], device='cuda:0')



e 0, step 2000, loss 18.759445190429688

Target: tensor([5, 7, 7, 6, 4, 3, 1, 5], device='cuda:0')

Pred: tensor([5, 7, 7, 6, 4, 0, 1, 5], device='cuda:0')

5 7 7 6 4 3 1 5

e 1, step 4000, loss 13.546051025390625

Target: tensor([2, 2, 6, 1, 4, 0, 3, 6], device='cuda:0')

Pred: tensor([2, 2, 6, 1, 6, 0, 3, 6], device='cuda:0')

2 2 6 1 4 0 3 6

e 2, step 6000, loss 9.396353721618652

Target: tensor([3, 3, 1, 8, 7, 4, 6, 8], device='cuda:0')

Pred: tensor([3, 3, 1, 8, 7, 4, 6, 8], device='cuda:0')

3 3 1 8 7 4 6 8

e 3, step 8000, loss 6.7225799560546875

Target: tensor([7, 7, 0, 4, 5, 7, 6, 7], device='cuda:0')

Pred: tensor([7, 7, 0, 4, 5, 7, 6, 7], device='cuda:0')

7 7 0 4 5 7 6 7

e 4, step 10000, loss 4.936184883117676

Target: tensor([4, 8, 3, 7, 3, 2, 5, 2], device='cuda:0')

Pred: tensor([4, 8, 3, 7, 3, 2, 5, 2], device='cuda:0')

4 8 3 7 3 2 5 2

e 4, step 12000, loss 3.538346767425537

Target: tensor([1, 5, 5, 3, 5, 7, 4, 9], device='cuda:0')

Pred: tensor([1, 5, 5, 3, 5, 7, 4, 9], device='cuda:0')

1 5 5 3 5 7 4 9

e 5, step 14000, loss 2.6060330867767334

Target: tensor([1, 7, 8, 1, 1, 4, 1, 7], device='cuda:0')

Pred: tensor([1, 7, 8, 1, 1, 4, 1, 7], device='cuda:0')

1 7 8 1 1 4 1 7

e 6, step 16000, loss 2.0112721920013428

Target: tensor([1, 5, 2, 2, 4, 2, 2, 0], device='cuda:0')

Pred: tensor([1, 5, 2, 2, 4, 2, 2, 0], device='cuda:0')

1 5 2 2 4 2 2 0

e 7, step 18000, loss 1.5799720287322998

Target: tensor([7, 6, 8, 0, 8, 8, 2, 2], device='cuda:0')

Pred: tensor([7, 6, 8, 0, 8, 8, 2, 2], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 7, 6, 8, 0, 8, 8, 2, and 2.

e 8, step 20000, loss 1.06905996799469

Target: tensor([0, 4, 1, 3, 2, 1, 8, 1], device='cuda:0')

Pred: tensor([0, 4, 1, 3, 2, 1, 8, 1], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 0, 4, 1, 3, 2, 1, 8, and 1.

e 8, step 22000, loss 0.8557407855987549

Target: tensor([3, 4, 1, 6, 1, 4, 2, 0], device='cuda:0')

Pred: tensor([3, 4, 1, 6, 1, 4, 2, 0], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 3, 4, 1, 6, 1, 4, 2, and 0.

e 9, step 24000, loss 0.6700837016105652

Target: tensor([1, 1, 8, 0, 0, 1, 6, 0], device='cuda:0')

Pred: tensor([1, 1, 8, 0, 0, 1, 6, 0], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 1, 1, 8, 0, 0, 1, 6, and 0.

e 10, step 26000, loss 0.6538465023040771

Target: tensor([3, 0, 4, 8, 6, 5, 7, 2], device='cuda:0')

Pred: tensor([3, 2, 4, 8, 6, 5, 7, 2], device='cuda:0')

3 0 4 8 6 5 7 2

e 11, step 28000, loss 0.49029141664505005

Target: tensor([1, 6, 4, 0, 9, 6, 1, 3], device='cuda:0')

Pred: tensor([1, 6, 4, 0, 9, 6, 1, 3], device='cuda:0')

1 6 4 0 9 6 1 3

e 12, step 30000, loss 0.4289899170398712

Target: tensor([5, 9, 3, 4, 8, 1, 8, 7], device='cuda:0')

Pred: tensor([5, 9, 3, 4, 8, 1, 8, 9], device='cuda:0')

5 9 3 4 8 1 8 9

e 12, step 32000, loss 0.5417484045028687

Target: tensor([0, 2, 5, 5, 4, 0, 1, 6], device='cuda:0')

Pred: tensor([0, 2, 5, 5, 4, 0, 1, 6], device='cuda:0')

0 2 5 5 4 0 1 6

e 13, step 34000, loss 0.39304643869400024

Target: tensor([7, 5, 4, 9, 8, 0, 7, 1], device='cuda:0')

Pred: tensor([7, 5, 4, 9, 8, 0, 7, 1], device='cuda:0')

7 5 4 9 8 0 7 1

e 14, step 36000, loss 0.2979143559932709

Target: tensor([3, 0, 2, 4, 1, 4, 3, 7], device='cuda:0')

Pred: tensor([3, 0, 2, 4, 1, 4, 3, 7], device='cuda:0')

3 0 2 4 1 4 3 7

e 15, step 38000, loss 0.42486128211021423

Target: tensor([2, 1, 7, 0, 1, 2, 8, 8], device='cuda:0')

Pred: tensor([3, 1, 7, 0, 1, 2, 8, 8], device='cuda:0')

2 1 7 0 1 2 8 8

e 16, step 40000, loss 0.2586093544960022

Target: tensor([7, 7, 5, 0, 1, 3, 1, 7], device='cuda:0')

Pred: tensor([7, 7, 5, 0, 1, 3, 1, 7], device='cuda:0')

7 7 5 0 1 3 1 7

e 16, step 42000, loss 0.5022361874580383

Target: tensor([6, 9, 9, 4, 8, 9, 0, 7], device='cuda:0')

Pred: tensor([6, 3, 9, 4, 8, 9, 0, 7], device='cuda:0')

6 9 9 4 8 9 0 7

e 17, step 44000, loss 0.2821335196495056

Target: tensor([7, 8, 2, 7, 8, 8, 7, 5], device='cuda:0')

Pred: tensor([7, 8, 2, 7, 8, 8, 7, 5], device='cuda:0')

7 8 2 7 8 8 7 5

e 18, step 46000, loss 0.2434415966272354

Target: tensor([9, 2, 1, 1, 3, 4, 1, 7], device='cuda:0')

Pred: tensor([9, 2, 1, 1, 3, 4, 1, 7], device='cuda:0')

9 2 1 1 3 4 1 7

e 19, step 48000, loss 0.2974969744682312

Target: tensor([8, 2, 5, 4, 8, 3, 0, 7], device='cuda:0')

Pred: tensor([8, 2, 5, 4, 8, 3, 0, 7], device='cuda:0')

8 2 5 4 8 3 0 7

e 20, step 50000, loss 0.251396507024765

Target: tensor([7, 2, 9, 6, 2, 1, 6, 7], device='cuda:0')

Pred: tensor([7, 2, 9, 6, 2, 1, 6, 7], device='cuda:0')

7 2 9 6 2 1 6 7

e 20, step 52000, loss 0.2520684003829956

Target: tensor([1, 8, 1, 7, 6, 3, 1, 1], device='cuda:0')

Pred: tensor([1, 8, 1, 7, 6, 3, 1, 1], device='cuda:0')

1 8 1 7 6 3 1 1

e 21, step 54000, loss 0.24324554204940796

Target: tensor([6, 5, 1, 9, 9, 9, 6, 7], device='cuda:0')

Pred: tensor([6, 5, 1, 9, 9, 9, 6, 7], device='cuda:0')

6 5 1 9 9 9 6 7

e 22, step 56000, loss 0.2629275619983673

Target: tensor([9, 2, 8, 8, 9, 4, 8, 0], device='cuda:0')

Pred: tensor([9, 2, 8, 8, 9, 4, 8, 0], device='cuda:0')

9 2 8 8 4 4 8 0

e 23, step 58000, loss 0.2249755859375

Target: tensor([9, 6, 0, 8, 4, 7, 9, 7], device='cuda:0')

Pred: tensor([9, 6, 0, 8, 4, 7, 9, 7], device='cuda:0')

9 6 0 8 4 7 9 7

e 24, step 60000, loss 0.28642648458480835

Target: tensor([7, 1, 3, 8, 3, 7, 2, 2], device='cuda:0')

Pred: tensor([7, 1, 3, 8, 3, 7, 2, 2], device='cuda:0')

7 1 3 8 3 4 2 2

e 24, step 62000, loss 0.21237778663635254

Target: tensor([6, 8, 7, 3, 9, 7, 8, 3], device='cuda:0')

Pred: tensor([6, 8, 7, 3, 9, 7, 8, 3], device='cuda:0')

6 8 7 3 9 7 8 3

e 25, step 64000, loss 0.3239150941371918

Target: tensor([3, 5, 0, 4, 1, 7, 4, 8], device='cuda:0')

Pred: tensor([3, 5, 0, 4, 1, 7, 4, 8], device='cuda:0')

3 5 0 4 1 7 4 8

e 26, step 66000, loss 0.20264635980129242

Target: tensor([7, 4, 6, 6, 2, 2, 9, 0], device='cuda:0')

Pred: tensor([7, 4, 6, 6, 2, 2, 9, 0], device='cuda:0')

7 4 6 6 2 2 9 0

e 27, step 68000, loss 0.25085484981536865

Target: tensor([9, 3, 1, 0, 2, 6, 7, 4], device='cuda:0')

Pred: tensor([9, 3, 1, 0, 2, 6, 7, 4], device='cuda:0')

9 3 1 0 2 6 7 4

e 28, step 70000, loss 0.34856224060058594

Target: tensor([3, 9, 8, 5, 7, 1, 6, 9], device='cuda:0')

Pred: tensor([3, 9, 8, 5, 7, 1, 6, 9], device='cuda:0')

3 9 8 5 7 1 6 9

e 28, step 72000, loss 0.27126139402389526

Target: tensor([6, 0, 5, 3, 4, 0, 2, 0], device='cuda:0')

Pred: tensor([6, 0, 5, 3, 4, 0, 2, 0], device='cuda:0')

6 0 5 3 4 0 2 0

e 29, step 74000, loss 0.3506920337677002

Target: tensor([9, 7, 7, 7, 2, 5, 4, 0], device='cuda:0')

Pred: tensor([9, 7, 7, 7, 2, 5, 4, 0], device='cuda:0')

9 7 7 7 2 5 4 0

e 30, step 76000, loss 0.20774713158607483

Target: tensor([7, 4, 4, 6, 0, 2, 0, 5], device='cuda:0')

Pred: tensor([7, 4, 4, 6, 0, 2, 0, 5], device='cuda:0')

7 4 4 6 0 2 0 5

e 31, step 78000, loss 0.22745007276535034

Target: tensor([7, 3, 9, 6, 6, 5, 6, 2], device='cuda:0')

Pred: tensor([7, 3, 9, 6, 6, 5, 6, 2], device='cuda:0')

7 3 9 6 6 5 6 2

e 32, step 80000, loss 0.23046478629112244

Target: tensor([7, 4, 6, 2, 3, 3, 8, 8], device='cuda:0')

Pred: tensor([7, 4, 6, 2, 3, 3, 8, 8], device='cuda:0')

7 4 6 2 3 3 8 8

e 32, step 82000, loss 0.20866142213344574

Target: tensor([7, 0, 8, 1, 8, 2, 8, 3], device='cuda:0')

Pred: tensor([7, 0, 8, 1, 8, 2, 8, 3], device='cuda:0')

7 0 8 1 8 2 8 3

e 33, step 84000, loss 0.26251378655433655

Target: tensor([3, 4, 8, 2, 3, 6, 2, 6], device='cuda:0')

Pred: tensor([3, 4, 8, 2, 3, 6, 2, 6], device='cuda:0')

3 4 8 2 3 6 2 6

e 34, step 86000, loss 0.26466721296310425

Target: tensor([6, 1, 7, 4, 2, 8, 7, 2], device='cuda:0')

Pred: tensor([6, 1, 7, 4, 2, 8, 7, 2], device='cuda:0')

6 1 7 4 2 8 7 2

e 35, step 88000, loss 0.2130233347415924

Target: tensor([8, 2, 0, 6, 8, 4, 0, 6], device='cuda:0')

Pred: tensor([8, 2, 0, 6, 8, 4, 0, 6], device='cuda:0')

8 2 0 6 8 4 0 6

e 36, step 90000, loss 0.26777076721191406

Target: tensor([7, 0, 2, 4, 9, 1, 0, 7], device='cuda:0')

Pred: tensor([7, 2, 2, 4, 9, 1, 0, 7], device='cuda:0')

7 0 2 4 9 1 0 7

e 36, step 92000, loss 0.2653909921646118
Target: tensor([4, 7, 2, 1, 0, 4, 5, 8], device='cuda:0')
Pred: tensor([4, 7, 2, 1, 0, 4, 5, 8], device='cuda:0')

4 7 2 1 0 4 5 8

e 37, step 94000, loss 0.20144307613372803
Target: tensor([4, 5, 0, 7, 4, 3, 7, 0], device='cuda:0')
Pred: tensor([4, 5, 0, 7, 4, 3, 7, 0], device='cuda:0')

4 5 0 7 4 3 7 0

e 38, step 96000, loss 0.3582932949066162
Target: tensor([3, 3, 1, 8, 8, 4, 6, 2], device='cuda:0')
Pred: tensor([3, 3, 1, 8, 8, 4, 6, 2], device='cuda:0')

3 3 1 8 8 4 6 2

e 39, step 98000, loss 0.20364491641521454
Target: tensor([0, 1, 3, 4, 6, 1, 7, 4], device='cuda:0')
Pred: tensor([0, 1, 3, 4, 6, 1, 7, 4], device='cuda:0')

0 1 3 4 6 1 7 4

e 40, step 100000, loss 0.21123895049095154
Target: tensor([7, 5, 6, 8, 6, 9, 0, 3], device='cuda:0')
Pred: tensor([7, 5, 6, 8, 6, 9, 0, 3], device='cuda:0')

7 5 6 8 6 9 0 3

e 40, step 102000, loss 0.30827152729034424
Target: tensor([2, 1, 2, 7, 6, 4, 0, 2], device='cuda:0')
Pred: tensor([2, 1, 2, 7, 6, 4, 0, 2], device='cuda:0')

2 1 2 7 6 4 0 2

e 41, step 104000, loss 0.40713948011398315
Target: tensor([3, 0, 3, 4, 2, 5, 4, 3], device='cuda:0')
Pred: tensor([3, 0, 3, 4, 2, 5, 4, 3], device='cuda:0')

3 0 3 4 2 5 4 3

e 42, step 106000, loss 0.22978179156780243
Target: tensor([4, 2, 8, 6, 3, 9, 7, 1], device='cuda:0')
Pred: tensor([4, 2, 8, 6, 3, 9, 7, 1], device='cuda:0')

4 2 8 6 3 9 7 1

e 43, step 108000, loss 0.43083497881889343

Target: tensor([8, 3, 2, 0, 3, 7, 5, 7], device='cuda:0')

Pred: tensor([8, 3, 2, 0, 3, 7, 5, 2], device='cuda:0')

8 3 2 0 3 7 5 7

e 44, step 110000, loss 0.24697725474834442

Target: tensor([7, 5, 5, 5, 8, 4, 9, 2], device='cuda:0')

Pred: tensor([7, 5, 5, 5, 8, 4, 9, 2], device='cuda:0')

7 5 5 5 8 4 9 2

e 44, step 112000, loss 0.24559302628040314

Target: tensor([6, 5, 3, 3, 1, 9, 9, 1], device='cuda:0')

Pred: tensor([6, 5, 3, 3, 1, 4, 9, 1], device='cuda:0')

6 5 3 3 1 9 9 1

e 45, step 114000, loss 0.22144173085689545

Target: tensor([7, 8, 5, 8, 0, 6, 4, 9], device='cuda:0')

Pred: tensor([7, 8, 5, 8, 0, 6, 4, 9], device='cuda:0')

7 8 5 8 0 6 4 9

e 46, step 116000, loss 0.1986529380083084

Target: tensor([5, 7, 9, 5, 3, 2, 8, 3], device='cuda:0')

Pred: tensor([5, 7, 9, 5, 3, 2, 8, 3], device='cuda:0')

5 7 9 5 3 2 8 3

e 47, step 118000, loss 0.2182689905166626

Target: tensor([8, 5, 7, 8, 9, 7, 2, 5], device='cuda:0')

Pred: tensor([8, 5, 7, 8, 9, 7, 2, 5], device='cuda:0')

8 5 7 8 9 7 2 5

e 48, step 120000, loss 0.24669942259788513

Target: tensor([6, 6, 3, 8, 2, 5, 7, 8], device='cuda:0')

Pred: tensor([6, 6, 3, 8, 2, 5, 7, 8], device='cuda:0')

6 6 3 8 2 5 7 8

e 48, step 122000, loss 0.21693159639835358

Target: tensor([1, 4, 7, 5, 8, 7, 5, 8], device='cuda:0')

Pred: tensor([1, 4, 7, 5, 8, 7, 5, 8], device='cuda:0')

1 4 7 5 8 7 5 8

e 49, step 124000, loss 0.21607306599617004

Target: tensor([6, 1, 0, 5, 3, 9, 6, 9], device='cuda:0')

Pred: tensor([6, 1, 0, 5, 3, 9, 6, 9], device='cuda:0')

6 1 0 5 3 9 6 9

e 50, step 126000, loss 0.30115222930908203

Target: tensor([0, 3, 9, 6, 5, 9, 1, 7], device='cuda:0')

Pred: tensor([0, 3, 9, 6, 5, 9, 1, 7], device='cuda:0')

0 3 9 6 5 9 1 7

e 51, step 128000, loss 0.34286022186279297

Target: tensor([0, 9, 8, 3, 6, 3, 1, 4], device='cuda:0')

Pred: tensor([0, 9, 8, 3, 6, 2, 1, 4], device='cuda:0')

0 9 8 3 6 2 1 4

e 52, step 130000, loss 0.3344825506210327

Target: tensor([7, 8, 1, 4, 4, 2, 6, 4], device='cuda:0')

Pred: tensor([4, 8, 1, 9, 4, 2, 6, 4], device='cuda:0')

→ 8 1 4 4 2 6 4

e 52, step 132000, loss 0.22589246928691864

Target: tensor([6, 2, 7, 5, 6, 8, 1, 1], device='cuda:0')

Pred: tensor([6, 2, 7, 5, 6, 8, 1, 1], device='cuda:0')

6 2 7 5 6 8 1 1

e 53, step 134000, loss 0.20860764384269714

Target: tensor([6, 5, 0, 9, 9, 8, 5, 0], device='cuda:0')

Pred: tensor([6, 5, 0, 9, 9, 8, 5, 0], device='cuda:0')

6 5 0 9 9 8 5 0

e 54, step 136000, loss 0.1999923586845398

Target: tensor([2, 2, 9, 4, 2, 4, 2, 9], device='cuda:0')

Pred: tensor([2, 2, 9, 4, 2, 4, 2, 9], device='cuda:0')

2 2 9 4 2 4 2 9

e 55, step 138000, loss 0.19996346533298492

Target: tensor([3, 1, 7, 6, 2, 0, 2, 1], device='cuda:0')

Pred: tensor([3, 1, 7, 6, 2, 0, 2, 1], device='cuda:0')

3 1 7 6 2 0 2 1

e 56, step 140000, loss 0.19718068838119507

Target: tensor([1, 8, 1, 1, 5, 7, 6, 3], device='cuda:0')

Pred: tensor([1, 8, 1, 1, 5, 7, 6, 3], device='cuda:0')

1 8 1 1 5 7 6 3

e 56, step 142000, loss 0.2202605903148651

Target: tensor([3, 4, 0, 9, 7, 5, 0, 0], device='cuda:0')

Pred: tensor([3, 4, 0, 9, 7, 5, 0, 0], device='cuda:0')

3 4 0 9 7 5 0 0

e 57, step 144000, loss 0.23822985589504242

Target: tensor([6, 2, 9, 0, 9, 7, 6, 5], device='cuda:0')

Pred: tensor([6, 2, 9, 0, 9, 7, 0, 5], device='cuda:0')

6 2 9 0 9 7 6 5

e 58, step 146000, loss 0.18683098256587982

Target: tensor([9, 4, 4, 3, 6, 4, 9, 7], device='cuda:0')

Pred: tensor([9, 4, 4, 3, 6, 4, 9, 7], device='cuda:0')

9 4 4 3 6 4 9 7

e 59, step 148000, loss 0.4197006821632385

Target: tensor([3, 2, 2, 3, 3, 5, 8, 4], device='cuda:0')

Pred: tensor([3, 2, 3, 3, 3, 5, 8, 4], device='cuda:0')

3 2 2 3 3 5 8 4

e 60, step 150000, loss 0.2731650471687317

Target: tensor([0, 6, 2, 7, 7, 5, 6, 1], device='cuda:0')

Pred: tensor([0, 6, 2, 7, 7, 5, 6, 8], device='cuda:0')

0 6 2 7 7 5 6 1

e 60, step 152000, loss 0.2503570020198822

Target: tensor([7, 3, 2, 1, 1, 8, 7, 1], device='cuda:0')

Pred: tensor([7, 3, 2, 1, 1, 8, 7, 1], device='cuda:0')

7 3 2 1 1 8 7 1

e 61, step 154000, loss 0.20371848344802856

Target: tensor([3, 4, 7, 0, 8, 7, 0, 3], device='cuda:0')

Pred: tensor([3, 4, 7, 0, 8, 7, 0, 3], device='cuda:0')

3 4 7 0 8 7 0 3

e 62, step 156000, loss 0.21197852492332458

Target: tensor([7, 0, 8, 5, 1, 9, 9, 5], device='cuda:0')

Pred: tensor([7, 0, 8, 5, 1, 9, 9, 5], device='cuda:0')

7 0 8 5 1 9 9 5

e 63, step 158000, loss 0.24664781987667084

Target: tensor([3, 0, 2, 4, 9, 7, 7, 1], device='cuda:0')

Pred: tensor([3, 0, 2, 4, 9, 7, 7, 1], device='cuda:0')

3 0 2 4 9 7 7 1

e 64, step 160000, loss 0.23666898906230927

Target: tensor([9, 5, 0, 2, 8, 1, 7, 2], device='cuda:0')

Pred: tensor([9, 5, 0, 2, 8, 1, 7, 2], device='cuda:0')

9 5 0 2 8 1 7 2

e 64, step 162000, loss 0.21137672662734985

Target: tensor([6, 4, 6, 8, 7, 6, 5, 7], device='cuda:0')

Pred: tensor([6, 4, 6, 8, 7, 6, 5, 7], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 6, 4, 6, 8, 7, 6, 5, and 7, written in a slightly cursive style.

e 65, step 164000, loss 0.25137513875961304
Target: tensor([4, 1, 9, 5, 9, 6, 8, 0], device='cuda:0')
Pred: tensor([4, 1, 9, 5, 9, 6, 8, 0], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 4, 1, 9, 5, 9, 6, 8, and 0, written in a slightly cursive style.

e 66, step 166000, loss 0.21092897653579712
Target: tensor([5, 5, 8, 3, 4, 1, 5, 6], device='cuda:0')
Pred: tensor([5, 5, 8, 3, 4, 1, 5, 6], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 5, 5, 8, 3, 4, 1, 5, and 6, written in a slightly cursive style.

e 67, step 168000, loss 0.41123777627944946
Target: tensor([9, 4, 5, 0, 3, 8, 8, 2], device='cuda:0')
Pred: tensor([9, 4, 5, 0, 9, 8, 8, 2], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 9, 4, 5, 0, 3, 8, 8, and 2, written in a slightly cursive style.

e 68, step 170000, loss 0.2300581932067871
Target: tensor([6, 9, 1, 3, 3, 0, 3, 9], device='cuda:0')
Pred: tensor([6, 9, 1, 3, 3, 0, 3, 9], device='cuda:0')

6 9 / 3 3 0 3 7

e 68, step 172000, loss 0.4018319547176361

Target: tensor([1, 1, 9, 0, 0, 7, 4, 0], device='cuda:0')

Pred: tensor([1, 1, 9, 0, 0, 7, 4, 6], device='cuda:0')

1 / 9 0 0 7 4 6

e 69, step 174000, loss 0.3300042152404785

Target: tensor([4, 0, 5, 6, 8, 6, 2, 5], device='cuda:0')

Pred: tensor([4, 0, 5, 6, 8, 6, 2, 5], device='cuda:0')

4 0 5 6 8 6 2 5

e 70, step 176000, loss 0.31812232732772827

Target: tensor([4, 9, 2, 8, 7, 0, 1, 1], device='cuda:0')

Pred: tensor([4, 9, 2, 8, 7, 0, 1, 1], device='cuda:0')

4 9 2 8 7 0 1 1

e 71, step 178000, loss 0.23327434062957764

Target: tensor([9, 1, 0, 5, 3, 8, 3, 4], device='cuda:0')

Pred: tensor([9, 1, 0, 5, 3, 8, 3, 4], device='cuda:0')

9 1 0 5 3 8 3 4

e 72, step 180000, loss 0.25247544050216675

Target: tensor([8, 4, 2, 1, 1, 1, 5, 4], device='cuda:0')

Pred: tensor([8, 4, 2, 1, 1, 1, 5, 4], device='cuda:0')

8 4 2 1 1 1 5 4

e 72, step 182000, loss 0.39648184180259705

Target: tensor([5, 5, 6, 1, 3, 2, 1, 4], device='cuda:0')

Pred: tensor([5, 5, 6, 1, 3, 2, 1, 4], device='cuda:0')

5 5 6 1 3 2 1 4

e 73, step 184000, loss 0.4915977120399475

Target: tensor([7, 9, 4, 0, 8, 6, 7, 7], device='cuda:0')

Pred: tensor([4, 9, 4, 0, 8, 6, 7, 7], device='cuda:0')

7 9 4 0 8 6 7 7

e 74, step 186000, loss 0.19474978744983673

Target: tensor([6, 8, 2, 6, 7, 8, 7, 4], device='cuda:0')

Pred: tensor([6, 8, 2, 6, 7, 8, 7, 4], device='cuda:0')

6 8 2 6 7 8 7 4

e 75, step 188000, loss 0.5243025422096252

Target: tensor([1, 8, 4, 9, 3, 3, 8, 3], device='cuda:0')

Pred: tensor([1, 8, 4, 9, 3, 3, 8, 3], device='cuda:0')

1 8 4 9 3 3 8 3

e 76, step 190000, loss 0.20132921636104584

Target: tensor([9, 8, 5, 2, 2, 6, 5, 4], device='cuda:0')

Pred: tensor([9, 8, 5, 2, 2, 6, 5, 4], device='cuda:0')

9 8 5 2 2 6 5 4

e 76, step 192000, loss 0.24694707989692688

Target: tensor([2, 9, 7, 0, 4, 0, 0, 2], device='cuda:0')

Pred: tensor([2, 9, 7, 0, 4, 0, 0, 2], device='cuda:0')

2 9 7 0 4 0 0 2

e 77, step 194000, loss 0.1933869570493698

Target: tensor([7, 8, 0, 3, 2, 1, 7, 6], device='cuda:0')

Pred: tensor([7, 8, 0, 3, 2, 1, 7, 6], device='cuda:0')

7 8 0 3 2 / 7 6

e 78, step 196000, loss 0.2940853238105774

Target: tensor([2, 3, 7, 2, 6, 8, 3, 6], device='cuda:0')

Pred: tensor([8, 3, 7, 2, 6, 8, 3, 6], device='cuda:0')

2 3 7 2 6 8 3 6

e 79, step 198000, loss 0.237261563539505

Target: tensor([8, 9, 6, 7, 3, 4, 2, 5], device='cuda:0')

Pred: tensor([8, 9, 6, 7, 3, 4, 2, 5], device='cuda:0')

8 9 6 7 3 4 2 5

e 80, step 200000, loss 0.19705039262771606

Target: tensor([4, 3, 2, 2, 6, 1, 3, 6], device='cuda:0')

Pred: tensor([4, 3, 2, 2, 6, 1, 3, 6], device='cuda:0')

4 3 2 2 6 1 3 6

e 80, step 202000, loss 0.444485068321228

Target: tensor([9, 1, 7, 3, 2, 6, 4, 8], device='cuda:0')

Pred: tensor([9, 1, 7, 3, 2, 6, 4, 8], device='cuda:0')

9 1 7 3 2 6 4 8

e 81, step 204000, loss 0.402957022190094

Target: tensor([0, 7, 5, 1, 1, 5, 0, 8], device='cuda:0')

Pred: tensor([0, 7, 9, 1, 1, 5, 0, 8], device='cuda:0')

0 7 5 1 1 5 0 8

e 82, step 206000, loss 0.26242944598197937

Target: tensor([3, 5, 7, 0, 6, 1, 9, 6], device='cuda:0')

Pred: tensor([3, 5, 7, 0, 6, 1, 9, 6], device='cuda:0')

3 5 7 0 6 1 9 6

e 83, step 208000, loss 0.25037142634391785

Target: tensor([5, 6, 3, 8, 2, 3, 1, 1], device='cuda:0')

Pred: tensor([5, 6, 3, 8, 2, 3, 1, 1], device='cuda:0')

5 6 3 8 2 3 1 1

e 84, step 210000, loss 0.2539176344871521

Target: tensor([1, 6, 0, 5, 8, 7, 1, 9], device='cuda:0')

Pred: tensor([1, 6, 0, 5, 8, 7, 1, 9], device='cuda:0')

1 6 0 5 8 7 1 9

e 84, step 212000, loss 0.3109821677207947

Target: tensor([3, 3, 2, 1, 3, 0, 6, 7], device='cuda:0')

Pred: tensor([3, 3, 2, 1, 3, 0, 6, 7], device='cuda:0')

3 3 2 1 3 0 6 7

e 85, step 214000, loss 0.3416498005390167

Target: tensor([0, 8, 8, 1, 4, 9, 9, 1], device='cuda:0')

Pred: tensor([0, 8, 8, 1, 4, 9, 9, 1], device='cuda:0')

0 8 8 1 4 9 9 1

e 86, step 216000, loss 0.2438945472240448

Target: tensor([8, 2, 2, 2, 8, 4, 1, 4], device='cuda:0')

Pred: tensor([8, 2, 2, 2, 8, 4, 1, 4], device='cuda:0')

8 2 2 2 8 4 1 4

e 87, step 218000, loss 0.3577214479446411

Target: tensor([3, 4, 3, 9, 8, 9, 0, 0], device='cuda:0')

Pred: tensor([3, 1, 3, 9, 8, 9, 0, 0], device='cuda:0')

3 4 3 9 8 9 0 0

e 88, step 220000, loss 0.2549368441104889

Target: tensor([7, 4, 9, 6, 4, 5, 8, 7], device='cuda:0')

Pred: tensor([7, 4, 9, 6, 4, 5, 8, 7], device='cuda:0')

7 4 9 6 4 5 8 7

e 88, step 222000, loss 0.20837688446044922

Target: tensor([4, 2, 9, 1, 1, 4, 1, 5], device='cuda:0')

Pred: tensor([4, 2, 9, 1, 1, 4, 1, 5], device='cuda:0')

4 2 9 \ / 4 1 5

e 89, step 224000, loss 0.2294844686985016

Target: tensor([7, 0, 1, 1, 1, 8, 1, 3], device='cuda:0')

Pred: tensor([7, 0, 1, 1, 1, 8, 1, 3], device='cuda:0')

7 0 1 1 1 8 1 3

e 90, step 226000, loss 0.4892371892929077

Target: tensor([5, 0, 4, 7, 8, 8, 6, 6], device='cuda:0')

Pred: tensor([5, 0, 8, 7, 8, 8, 6, 6], device='cuda:0')

5 0 8 7 8 8 6 6

e 91, step 228000, loss 0.2625042796134949

Target: tensor([8, 0, 7, 2, 5, 4, 0, 5], device='cuda:0')

Pred: tensor([8, 0, 1, 2, 5, 4, 0, 5], device='cuda:0')

8 0 1 2 5 4 0 5

e 92, step 230000, loss 0.2328953891992569

Target: tensor([3, 1, 2, 4, 7, 2, 8, 2], device='cuda:0')

Pred: tensor([3, 1, 2, 4, 7, 2, 8, 2], device='cuda:0')

3 1 2 4 7 2 8 2

e 92, step 232000, loss 0.22881978750228882

Target: tensor([3, 3, 7, 0, 3, 1, 1, 1], device='cuda:0')

Pred: tensor([3, 3, 7, 0, 3, 1, 1, 1], device='cuda:0')

3 3 7 0 3 1 1 1

e 93, step 234000, loss 0.18785811960697174

Target: tensor([8, 6, 5, 7, 5, 8, 0, 9], device='cuda:0')

Pred: tensor([8, 6, 5, 7, 5, 8, 0, 9], device='cuda:0')

8 6 5 7 5 8 0 9

e 94, step 236000, loss 0.2069055736064911

Target: tensor([0, 5, 4, 2, 7, 6, 5, 6], device='cuda:0')

Pred: tensor([0, 5, 4, 2, 7, 6, 5, 6], device='cuda:0')

0 5 4 2 7 6 5 6

e 95, step 238000, loss 0.23149007558822632

Target: tensor([6, 0, 6, 1, 0, 5, 0, 3], device='cuda:0')

Pred: tensor([6, 0, 6, 1, 0, 5, 0, 3], device='cuda:0')

6 0 6 1 0 5 0 3

e 96, step 240000, loss 0.20913715660572052

Target: tensor([1, 2, 0, 0, 5, 0, 3, 5], device='cuda:0')

Pred: tensor([1, 2, 0, 0, 5, 0, 3, 5], device='cuda:0')

1 2 0 0 5 0 3 5

e 96, step 242000, loss 0.20399980247020721

Target: tensor([6, 6, 9, 7, 9, 6, 0, 2], device='cuda:0')

Pred: tensor([6, 6, 9, 7, 9, 6, 0, 2], device='cuda:0')

6 6 9 7 9 6 0 2

e 97, step 244000, loss 0.225369393825531

Target: tensor([0, 6, 6, 5, 7, 9, 1, 5], device='cuda:0')

Pred: tensor([0, 6, 6, 5, 7, 9, 1, 5], device='cuda:0')

0 6 6 5 7 9 1 5

e 98, step 246000, loss 0.24798017740249634

Target: tensor([3, 6, 0, 7, 6, 5, 0, 4], device='cuda:0')

Pred: tensor([3, 6, 0, 7, 6, 5, 0, 4], device='cuda:0')



3 6 0 7 6 5 0 4

e 99, step 248000, loss 0.19269657135009766

Target: tensor([2, 1, 3, 0, 0, 1, 7, 4], device='cuda:0')

Pred: tensor([2, 1, 3, 0, 0, 1, 7, 4], device='cuda:0')

2 1 3 0 0 1 7 4



1.14 Batch Normalization

[Batch Normalization](#) can accelerate the training procedure by shifting and scaling the input for each layer.

To-do: - (10 points) Complete the `batch_norm_forward()` function for class `BatchNorm`. - (5 points) Modify the class `LeNetDrop` to class `LeNetDropNorm`: add a batch normalization layer after layer each dropout layer for `LeNetDrop`. - Write the dataset loading script, training script for `LeNetDropNorm`. - Train `LeNetDropNorm` on the full MNIST dataset. - (10 points) Plot out the training loss curve, validation loss curve, training accuracy curve and validation accuracy curve.

```
[25]: class BatchNorm(nn.Module):
    def __init__(self, num_features, device, eps=1e-5):
        super(BatchNorm, self).__init__()
        self.gamma = nn.Parameter(torch.ones(num_features, dtype=torch.float32,
        ↪device=device), requires_grad=True)
        self.beta = nn.Parameter(torch.zeros(num_features, dtype=torch.float32,
        ↪device=device), requires_grad=True)
        self.moving_mean = 0.
        self.moving_var = 0.
        self.eps = eps
        self.num_features = num_features
        self.device = device

    def batch_norm_forward(self, x, training=True):
        """
        input:
            X: (N, *), input of dropout layer, where N is the batch size
        output:
            Y: (N, *), where N is the batch size
        """
        ### Start the code here ###
        if training:
            if len(x.shape) == 2:
                mean = x.mean(0).unsqueeze(0).expand_as(x).cuda()
                var = ((x-mean)**2).mean(0).unsqueeze(0).expand_as(x).cuda()
            else:
                mean = x.mean(-1).mean(-1).mean(0).unsqueeze(-1).unsqueeze(-1).
        ↪unsqueeze(0).expand_as(x).to(self.device)
                var = ((x-mean)**2).mean(-1).mean(-1).mean(0).unsqueeze(-1).
        ↪unsqueeze(-1).unsqueeze(0).expand_as(x).to(self.device)
                x_hat = (x-mean) / (var + self.eps)**(0.5)
                momentum = 0.95
                self.moving_mean = momentum*self.moving_mean+(1-momentum)*mean
                self.moving_var = momentum*self.moving_var+(1-momentum)*var
            else:
                x_hat = (x - self.moving_mean) / (self.moving_var + self.eps)**0.5

        Y = self.gamma*x_hat.cuda() + self.beta.cuda()
        ### End of the code ###

        return Y
```



```
def forward(self, inputs, training=True):
    return self.batch_norm_forward(inputs, training)
```

```
[26]: # dataset loading, training script and other codes starts here:
# LeNetDrop sketch code
class LeNetDropNorm(nn.Module):
    def __init__(self, img_c, device):
        super(LeNetDropNorm, self).__init__()
        drop_p = 0.2
        self.c1 = Conv2D(img_c, 6, [5,5], [1,1], [2,2,2,2], device)
        self.d1 = Dropout(drop_p, device)
        self.n1 = BatchNorm([1, 6, 1, 1], device)

        self.p2 = nn.MaxPool2d(2, stride=2)
        self.c3 = Conv2D(6, 16, [5,5], [1,1], [0,0,0,0], device)
        self.d3 = Dropout(drop_p, device)
        self.n3 = BatchNorm([1, 16, 1, 1], device)

        self.p4 = nn.MaxPool2d(2, stride=2)
        self.f5 = Linear(400, 120, device)
        self.d5 = Dropout(drop_p, device)
        self.n5 = BatchNorm([1,120], device)

        self.f6 = Linear(120, 84, device)
        self.f7 = Linear(84, 10, device)
        self.device = device

    def forward(self, imgs, labels, training=True):
        """
        inputs:
            imgs: (N, C, H, W), training samples from the MNIST training set, where N
            ↪ is the number of samples (batch_size),
            C is the image color channle number, H and W are the spatial size of
            ↪ the input images.
            labels: (N, L), ground truth for the input images, where N is the number
            ↪ of samples (batch_size) and L is the
            number of classes.
        outputs:
            loss: (1,), mean loss value over this batch of inputs.
        """
        N = imgs.shape[0]

        o_c1 = F.relu(self.c1(imgs))
        o_d1 = self.d1(o_c1, training)
        o_n1 = self.n1(o_d1, training)
```

```

o_p2 = self.p2(o_n1)
o_c3 = F.relu(self.c3(o_p2))
o_d3 = self.d3(o_c3, training)
o_n3 = self.n3(o_d3, training)

o_p4 = self.p4(o_n3)

### Start the code here ###
# 1. Please complete the rest of LeNet to get the scores predicted by LeNet
→for each input images #
# need to flatten the matrix before forwarding to the dense layer
o_f5 = F.relu(self.f5(o_p4.reshape(o_p4.shape[0], -1)))
o_d5 = self.d5(o_f5, training)
o_n5 = self.n5(o_d5, training)

o_f6 = F.relu(self.f6(o_n5))
o_f7 = self.f7(o_f6)
# 2. Please use the implemented objective function to obtain the losses of
→each input. #
p = softmax1d(o_f7)
# 3. We will return the mean value of the losses. #
loss = cross_entropy_loss(p, labels)
### End of the code ###
return loss.mean(), p

```

```

[27]: #####
# Train LeNetDrop on the full dataset
#####
# Dataset loading script
train_ds = MNISTDataset(60000, split = "training")
val_ds = MNISTDataset(10000, split = "validation")

# training batch size, hyper-parameter
batch_size = 24

# dataset loader
train_dl = DataLoader(train_ds, batch_size=batch_size, shuffle=True,
→drop_last=True)
val_dl = DataLoader(val_ds, batch_size=batch_size, shuffle=True, drop_last=True)
img_channel = 1

```

```

[28]: #####
# Training script
# learning rate, hyper-parameter
lr = 1e-3

```

```

# using GPU if it's available
# device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
device = torch.device('cuda')

img_channel = 1
net = LeNetDropNorm(img_channel, device)
# keep a list of model parameters
params = [p for (n, p) in net.named_parameters()]

# training epochs, hyper-parameter
epochs = 100

# keep tracking of the changing of loss and accuracy of predictions
train_loss_list = []
val_loss_list = []
train_acc_list = []
val_acc_list = []
print_interval = 2000
for e in range(epochs):
    net.train()
    for i, (imgs, lbs) in enumerate(train_dl):
        imgs = imgs.to(device)
        lbs = lbs.to(device)
        loss, prob = net(imgs, lbs)
        # add the weight decay penalty to the loss
        lamb = 2.0
        loss = weight_decay(loss, params, lamb, lbs.shape[0])

    net.zero_grad()
    grad = torch.autograd.grad(loss.to(device), params)

    # update weights
    lamb = 1
    step(params, grad, lr)

    # obtain the predictions
    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()
    if (i + e*len(train_dl)) % print_interval == 0:
        print("e {}, step {}, loss {}".format(e, i + e*len(train_dl), loss))
        print("Target:\t {} \nPred:\t {}".format(lbs[:8], pred[:8]))
        # visualize some samples
        imgs_to_vis = vutils.make_grid(imgs[:8].cpu()+125., nrow=8, pad_value=1)
        plt.imshow(imgs_to_vis.permute(1,2,0).numpy().astype(np.uint8))
        plt.axis("off")
        plt.show()
    train_loss_list.append(loss.detach().mean())

```

```

train_acc_list.append(acc.detach().mean())

net.eval()
for i, (imgs, lbs) in enumerate(val_dl):
    imgs = imgs.to(device)
    lbs = lbs.to(device)
    loss, prob = net(imgs, lbs, training=False)
    # add the weight decay penalty to the loss
    loss = weight_decay(loss, params, lamb, lbs.shape[0])

    pred = prob.argmax(dim=-1).view(batch_size)
    acc = (pred == lbs).float().mean()

val_loss_list.append(loss.detach().mean())
val_acc_list.append(acc.detach().mean())

# plotting logs
plt.plot(np.arange(epochs), train_loss_list, '-r',
         np.arange(epochs), val_loss_list, '-g')
plt.legend(('training error', 'validation error'))
plt.show()
plt.plot(np.arange(epochs), train_acc_list, '-r',
         np.arange(epochs), val_acc_list, '-g')
plt.legend(('training acc', 'validation acc'))
plt.show()

```

e 0, step 0, loss 34.09567642211914
Target: tensor([3, 3, 9, 2, 9, 2, 7, 4], device='cuda:0')
Pred: tensor([0, 5, 5, 6, 6, 6, 6, 5], device='cuda:0')




e 0, step 2000, loss 24.14310646057129
Target: tensor([0, 8, 1, 7, 6, 6, 9, 7], device='cuda:0')
Pred: tensor([0, 8, 1, 7, 6, 6, 9, 7], device='cuda:0')




e 1, step 4000, loss 17.499126434326172
Target: tensor([5, 6, 7, 7, 7, 8, 7, 5], device='cuda:0')
Pred: tensor([5, 6, 7, 7, 7, 3, 7, 9], device='cuda:0')



e 2, step 6000, loss 12.942313194274902
Target: tensor([4, 4, 1, 2, 2, 8, 0, 3], device='cuda:0')
Pred: tensor([4, 6, 1, 2, 2, 8, 0, 3], device='cuda:0')



e 3, step 8000, loss 9.888056755065918
Target: tensor([1, 3, 5, 4, 3, 0, 3, 8], device='cuda:0')
Pred: tensor([1, 5, 5, 4, 7, 0, 9, 8], device='cuda:0')



e 4, step 10000, loss 7.5098466873168945
Target: tensor([0, 8, 9, 8, 9, 5, 1, 5], device='cuda:0')
Pred: tensor([0, 8, 9, 8, 9, 5, 1, 5], device='cuda:0')



e 4, step 12000, loss 5.929827690124512

Target: tensor([1, 4, 2, 0, 7, 8, 7, 6], device='cuda:0')
Pred: tensor([1, 4, 2, 0, 7, 8, 7, 6], device='cuda:0')

A black rectangular box containing the handwritten digits '1 4 2 0 7 8 7 6' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 5, step 14000, loss 4.685754776000977
Target: tensor([1, 5, 3, 2, 2, 3, 8, 2], device='cuda:0')
Pred: tensor([1, 5, 3, 2, 2, 3, 8, 2], device='cuda:0')

A black rectangular box containing the handwritten digits '1 5 3 2 2 3 8 2' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 6, step 16000, loss 3.9934871196746826
Target: tensor([9, 7, 7, 6, 3, 8, 0, 8], device='cuda:0')
Pred: tensor([5, 7, 7, 6, 3, 8, 0, 8], device='cuda:0')

A black rectangular box containing the handwritten digits '9 7 7 6 3 8 0 8' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 7, step 18000, loss 3.1930806636810303
Target: tensor([0, 8, 7, 3, 3, 4, 3, 7], device='cuda:0')
Pred: tensor([0, 8, 7, 3, 3, 4, 3, 7], device='cuda:0')

A black rectangular box containing the handwritten digits '0 8 7 3 3 4 3 7' in white ink. The digits are slightly slanted and have a casual, handwritten style.

e 8, step 20000, loss 2.853916645050049
Target: tensor([8, 0, 2, 3, 7, 4, 6, 5], device='cuda:0')
Pred: tensor([8, 0, 2, 3, 7, 4, 6, 5], device='cuda:0')

8 0 2 3 7 4 6 5

e 8, step 22000, loss 2.6965689659118652
Target: tensor([4, 6, 0, 0, 6, 8, 2, 4], device='cuda:0')
Pred: tensor([4, 6, 0, 0, 6, 8, 2, 4], device='cuda:0')

4 6 0 0 6 8 2 4

e 9, step 24000, loss 2.5419158935546875
Target: tensor([9, 0, 4, 4, 1, 3, 9, 2], device='cuda:0')
Pred: tensor([9, 0, 4, 9, 1, 3, 9, 2], device='cuda:0')

9 0 4 4 1 3 9 2

e 10, step 26000, loss 2.250012159347534
Target: tensor([3, 3, 5, 3, 8, 8, 2, 0], device='cuda:0')
Pred: tensor([3, 3, 5, 3, 3, 8, 2, 0], device='cuda:0')

3 3 5 3 8 8 2 0

e 11, step 28000, loss 2.0988266468048096
Target: tensor([2, 3, 3, 5, 5, 9, 0, 5], device='cuda:0')
Pred: tensor([2, 3, 3, 5, 5, 9, 0, 5], device='cuda:0')

2 3 3 5 5 9 0 5

e 12, step 30000, loss 2.118147134780884

Target: tensor([4, 8, 9, 3, 0, 1, 2, 7], device='cuda:0')

Pred: tensor([4, 8, 9, 3, 5, 1, 2, 7], device='cuda:0')

4 8 9 3 6 1 2 7

e 12, step 32000, loss 1.9780573844909668

Target: tensor([8, 5, 4, 5, 2, 7, 3, 3], device='cuda:0')

Pred: tensor([8, 5, 4, 5, 2, 7, 3, 3], device='cuda:0')

8 5 4 5 2 7 3 3

e 13, step 34000, loss 1.923034906387329

Target: tensor([3, 0, 0, 5, 7, 0, 1, 4], device='cuda:0')

Pred: tensor([3, 0, 0, 5, 7, 0, 1, 4], device='cuda:0')

3 0 0 5 7 0 1 4

e 14, step 36000, loss 1.8998757600784302

Target: tensor([7, 6, 9, 9, 1, 2, 4, 2], device='cuda:0')

Pred: tensor([7, 6, 9, 9, 1, 2, 4, 2], device='cuda:0')

7 6 9 9 1 2 4 2

e 15, step 38000, loss 1.895589828491211
Target: tensor([8, 4, 7, 9, 8, 1, 8, 7], device='cuda:0')
Pred: tensor([8, 4, 7, 9, 8, 1, 8, 2], device='cuda:0')

8 4 7 9 8 1 8 7

e 16, step 40000, loss 1.861626148223877
Target: tensor([3, 2, 0, 4, 2, 6, 8, 9], device='cuda:0')
Pred: tensor([3, 2, 0, 4, 2, 6, 8, 9], device='cuda:0')

3 2 0 4 2 6 8 9

e 16, step 42000, loss 1.9671578407287598
Target: tensor([9, 6, 5, 3, 8, 8, 6, 5], device='cuda:0')
Pred: tensor([9, 6, 5, 3, 8, 8, 6, 5], device='cuda:0')

9 6 5 3 8 8 6 5

e 17, step 44000, loss 1.909401535987854
Target: tensor([2, 6, 5, 0, 0, 2, 9, 8], device='cuda:0')
Pred: tensor([2, 6, 5, 0, 0, 2, 9, 8], device='cuda:0')

26500298

e 18, step 46000, loss 1.8024046421051025

Target: tensor([6, 1, 1, 0, 0, 9, 9, 2], device='cuda:0')

Pred: tensor([6, 1, 1, 0, 0, 9, 9, 2], device='cuda:0')

61100992

e 19, step 48000, loss 1.8529415130615234

Target: tensor([9, 2, 6, 7, 3, 9, 5, 2], device='cuda:0')

Pred: tensor([9, 2, 6, 7, 3, 9, 5, 2], device='cuda:0')

92673952

e 20, step 50000, loss 1.8549593687057495

Target: tensor([4, 4, 1, 1, 9, 1, 3, 3], device='cuda:0')

Pred: tensor([4, 4, 1, 1, 9, 1, 3, 3], device='cuda:0')

44119133

e 20, step 52000, loss 1.8605475425720215

Target: tensor([1, 1, 7, 0, 2, 5, 8, 4], device='cuda:0')

Pred: tensor([1, 1, 7, 0, 2, 5, 8, 4], device='cuda:0')

1 1 7 0 2 5 8 4

e 21, step 54000, loss 1.8667407035827637

Target: tensor([1, 8, 4, 4, 8, 2, 3, 8], device='cuda:0')

Pred: tensor([1, 8, 4, 4, 8, 2, 3, 8], device='cuda:0')

1 8 4 4 8 2 3 8

e 22, step 56000, loss 1.976410984992981

Target: tensor([5, 6, 9, 8, 7, 8, 1, 4], device='cuda:0')

Pred: tensor([3, 6, 9, 8, 7, 8, 1, 4], device='cuda:0')

5 6 9 8 7 8 1 4

e 23, step 58000, loss 1.7720621824264526

Target: tensor([7, 8, 5, 6, 4, 7, 5, 8], device='cuda:0')

Pred: tensor([7, 8, 5, 6, 4, 7, 5, 8], device='cuda:0')

7 8 5 6 4 7 5 8

e 24, step 60000, loss 1.8417295217514038

Target: tensor([8, 3, 4, 5, 2, 1, 2, 0], device='cuda:0')

Pred: tensor([8, 3, 4, 5, 2, 1, 2, 0], device='cuda:0')

8 3 4 5 2 1 2 0

e 24, step 62000, loss 1.8653018474578857

Target: tensor([5, 5, 3, 4, 3, 5, 8, 4], device='cuda:0')

Pred: tensor([5, 5, 3, 4, 3, 5, 8, 4], device='cuda:0')

5 5 3 4 3 5 8 4

e 25, step 64000, loss 1.8613054752349854

Target: tensor([1, 4, 2, 1, 5, 9, 2, 6], device='cuda:0')

Pred: tensor([1, 9, 2, 1, 5, 9, 2, 6], device='cuda:0')

1 4 2 1 5 9 2 6

e 26, step 66000, loss 1.9290931224822998

Target: tensor([9, 5, 1, 0, 4, 9, 1, 0], device='cuda:0')

Pred: tensor([9, 5, 1, 0, 4, 9, 1, 0], device='cuda:0')

9 5 1 0 4 9 1 0

e 27, step 68000, loss 1.7754411697387695

Target: tensor([3, 8, 8, 3, 0, 0, 2, 2], device='cuda:0')

Pred: tensor([3, 8, 8, 3, 0, 0, 2, 2], device='cuda:0')

3 8 8 3 0 0 2 2

e 28, step 70000, loss 1.8111956119537354

Target: tensor([0, 2, 0, 4, 1, 4, 9, 1], device='cuda:0')

Pred: tensor([0, 2, 0, 4, 1, 4, 9, 1], device='cuda:0')

0 2 0 4 1 4 9 1

e 28, step 72000, loss 1.7836768627166748

Target: tensor([5, 3, 0, 2, 3, 4, 2, 1], device='cuda:0')

Pred: tensor([5, 3, 0, 2, 3, 4, 2, 1], device='cuda:0')

5 3 0 2 3 4 2 1

e 29, step 74000, loss 1.9323554039001465

Target: tensor([2, 6, 6, 2, 9, 2, 4, 5], device='cuda:0')

Pred: tensor([2, 6, 6, 2, 9, 2, 4, 5], device='cuda:0')

2 6 6 2 9 2 4 5

e 30, step 76000, loss 1.801836371421814

Target: tensor([3, 6, 9, 5, 3, 8, 3, 7], device='cuda:0')

Pred: tensor([3, 6, 9, 5, 3, 8, 3, 7], device='cuda:0')

3 6 9 5 3 8 3 7

e 31, step 78000, loss 1.8318942785263062

Target: tensor([9, 5, 4, 2, 2, 9, 0, 7], device='cuda:0')

Pred: tensor([9, 5, 4, 2, 2, 9, 0, 7], device='cuda:0')

9 5 4 2 2 9 0 7

e 32, step 80000, loss 1.8156256675720215

Target: tensor([9, 2, 0, 9, 1, 7, 6, 5], device='cuda:0')

Pred: tensor([9, 2, 0, 9, 1, 7, 6, 5], device='cuda:0')

9 2 0 9 1 7 6 5

e 32, step 82000, loss 1.8436275720596313

Target: tensor([5, 6, 3, 8, 1, 1, 4, 0], device='cuda:0')

Pred: tensor([5, 6, 3, 8, 2, 1, 9, 0], device='cuda:0')

5 6 3 8 1 1 4 0

e 33, step 84000, loss 1.8119051456451416

Target: tensor([3, 4, 7, 4, 1, 2, 6, 3], device='cuda:0')

Pred: tensor([3, 4, 7, 4, 1, 2, 6, 3], device='cuda:0')

3 4 7 4 1 2 6 3

e 34, step 86000, loss 1.8797708749771118
Target: tensor([2, 9, 3, 2, 5, 2, 6, 3], device='cuda:0')
Pred: tensor([2, 9, 3, 2, 5, 2, 6, 3], device='cuda:0')

2 9 3 2 5 2 6 3

e 35, step 88000, loss 1.8763177394866943
Target: tensor([3, 6, 2, 0, 5, 5, 4, 6], device='cuda:0')
Pred: tensor([3, 6, 2, 0, 5, 5, 4, 6], device='cuda:0')

3 6 2 0 5 5 4 6

e 36, step 90000, loss 1.8664355278015137
Target: tensor([3, 7, 2, 3, 4, 8, 8, 4], device='cuda:0')
Pred: tensor([3, 7, 2, 3, 4, 8, 8, 4], device='cuda:0')

3 7 2 3 4 8 8 4

e 36, step 92000, loss 1.802858829498291
Target: tensor([1, 4, 3, 6, 9, 5, 0, 5], device='cuda:0')
Pred: tensor([7, 4, 3, 6, 9, 5, 0, 5], device='cuda:0')

1 4 3 6 9 5 0 5

e 37, step 94000, loss 1.8296282291412354

Target: tensor([8, 5, 4, 5, 9, 0, 8, 2], device='cuda:0')

Pred: tensor([8, 5, 4, 5, 9, 0, 8, 2], device='cuda:0')

5 4 5 9 0 8 2

e 38, step 96000, loss 1.7463901042938232

Target: tensor([8, 0, 8, 5, 4, 1, 6, 9], device='cuda:0')

Pred: tensor([8, 0, 8, 5, 4, 1, 6, 9], device='cuda:0')

8 0 8 5 4 1 6 9

e 39, step 98000, loss 1.809220790863037

Target: tensor([3, 7, 4, 9, 6, 1, 3, 0], device='cuda:0')

Pred: tensor([3, 7, 4, 9, 6, 1, 3, 0], device='cuda:0')

3 7 4 9 6 1 3 0

e 40, step 100000, loss 1.855083703994751

Target: tensor([6, 0, 4, 2, 0, 4, 6, 0], device='cuda:0')

Pred: tensor([6, 0, 4, 2, 0, 4, 6, 0], device='cuda:0')

6 0 4 2 0 4 6 0

e 40, step 102000, loss 1.7687714099884033
Target: tensor([5, 6, 2, 1, 1, 2, 0, 4], device='cuda:0')
Pred: tensor([5, 6, 2, 1, 1, 2, 0, 4], device='cuda:0')

5 6 2 1 1 2 0 4

e 41, step 104000, loss 1.8525516986846924
Target: tensor([9, 1, 0, 0, 5, 6, 1, 0], device='cuda:0')
Pred: tensor([9, 1, 0, 0, 5, 6, 1, 0], device='cuda:0')

9 1 0 0 5 6 1 0

e 42, step 106000, loss 1.9518128633499146
Target: tensor([3, 9, 8, 1, 8, 3, 8, 0], device='cuda:0')
Pred: tensor([3, 9, 8, 1, 8, 3, 8, 0], device='cuda:0')

3 9 8 1 8 3 8 0

e 43, step 108000, loss 1.860915184020996
Target: tensor([1, 0, 2, 7, 1, 9, 3, 6], device='cuda:0')
Pred: tensor([1, 0, 2, 7, 1, 9, 3, 6], device='cuda:0')

1 0 2 7 1 9 3 6

e 44, step 110000, loss 1.7896907329559326

Target: tensor([1, 1, 4, 3, 2, 1, 7, 6], device='cuda:0')

Pred: tensor([1, 1, 4, 3, 2, 1, 7, 6], device='cuda:0')

1 1 4 3 2 1 7 6

e 44, step 112000, loss 1.9111802577972412

Target: tensor([3, 2, 7, 0, 1, 5, 8, 1], device='cuda:0')

Pred: tensor([3, 2, 7, 0, 1, 5, 8, 1], device='cuda:0')

3 2 7 0 1 5 8 1

e 45, step 114000, loss 1.831419825553894

Target: tensor([1, 0, 2, 4, 1, 0, 0, 3], device='cuda:0')

Pred: tensor([1, 0, 2, 4, 1, 0, 0, 3], device='cuda:0')

1 0 2 4 1 0 0 3

e 46, step 116000, loss 1.7968449592590332

Target: tensor([0, 7, 2, 9, 2, 3, 0, 7], device='cuda:0')

Pred: tensor([0, 7, 2, 9, 2, 3, 0, 7], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 0, 7, 2, 9, 2, 3, 0, and 7.

e 47, step 118000, loss 1.8440544605255127

Target: tensor([9, 0, 3, 1, 5, 1, 0, 0], device='cuda:0')

Pred: tensor([9, 0, 3, 1, 5, 1, 0, 0], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 9, 0, 3, 1, 5, 1, 0, and 0.

e 48, step 120000, loss 1.8538726568222046

Target: tensor([6, 5, 4, 4, 9, 7, 1, 5], device='cuda:0')

Pred: tensor([6, 8, 4, 4, 9, 7, 1, 5], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 6, 5, 4, 4, 9, 7, 1, and 5.

e 48, step 122000, loss 1.8428013324737549

Target: tensor([4, 2, 2, 9, 2, 4, 6, 3], device='cuda:0')

Pred: tensor([4, 2, 2, 9, 2, 4, 6, 3], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 4, 2, 2, 9, 2, 4, 6, and 3.

e 49, step 124000, loss 1.7825195789337158

Target: tensor([5, 6, 5, 0, 9, 7, 5, 2], device='cuda:0')

Pred: tensor([5, 6, 5, 0, 9, 7, 5, 2], device='cuda:0')

5 6 5 0 9 7 5 2

e 50, step 126000, loss 1.8047479391098022

Target: tensor([4, 0, 4, 3, 7, 8, 1, 4], device='cuda:0')

Pred: tensor([4, 0, 4, 3, 7, 8, 1, 4], device='cuda:0')

4 0 4 3 7 8 1 4

e 51, step 128000, loss 1.7868151664733887

Target: tensor([3, 9, 0, 3, 6, 5, 0, 8], device='cuda:0')

Pred: tensor([3, 9, 0, 3, 6, 5, 0, 8], device='cuda:0')

3 9 0 3 6 5 0 8

e 52, step 130000, loss 1.7983767986297607

Target: tensor([7, 9, 9, 3, 0, 8, 7, 0], device='cuda:0')

Pred: tensor([7, 9, 9, 3, 0, 8, 7, 0], device='cuda:0')

7 9 9 3 0 8 7 0

e 52, step 132000, loss 1.8953509330749512

Target: tensor([0, 5, 0, 1, 7, 3, 8, 2], device='cuda:0')

Pred: tensor([0, 5, 0, 1, 7, 3, 9, 2], device='cuda:0')

05017382

e 53, step 134000, loss 1.9458274841308594

Target: tensor([0, 1, 9, 6, 0, 5, 3, 4], device='cuda:0')

Pred: tensor([7, 1, 9, 6, 0, 5, 3, 4], device='cuda:0')

01.960534

e 54, step 136000, loss 1.8445227146148682

Target: tensor([8, 3, 5, 5, 5, 4, 3, 6], device='cuda:0')

Pred: tensor([8, 3, 5, 5, 5, 4, 3, 6], device='cuda:0')

83555436

e 55, step 138000, loss 1.7595453262329102

Target: tensor([1, 1, 3, 2, 1, 5, 5, 6], device='cuda:0')

Pred: tensor([1, 1, 3, 2, 1, 5, 5, 6], device='cuda:0')

11321556

e 56, step 140000, loss 1.8456008434295654

Target: tensor([7, 8, 8, 9, 2, 5, 5, 8], device='cuda:0')

Pred: tensor([7, 8, 8, 9, 2, 5, 5, 8], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 7, 8, 8, 1, 2, 5, 5, and 8.

e 56, step 142000, loss 1.7782552242279053

Target: tensor([8, 6, 0, 7, 1, 4, 0, 7], device='cuda:0')

Pred: tensor([8, 6, 0, 7, 1, 4, 0, 7], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 8, 6, 0, 7, 1, 4, 0, and 7.

e 57, step 144000, loss 1.7855638265609741

Target: tensor([5, 5, 2, 1, 2, 5, 6, 9], device='cuda:0')

Pred: tensor([5, 5, 2, 1, 2, 5, 6, 9], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 5, 5, 2, 1, 2, 5, 6, and 9.

e 58, step 146000, loss 1.7774146795272827

Target: tensor([8, 1, 9, 2, 7, 1, 8, 3], device='cuda:0')

Pred: tensor([8, 1, 9, 2, 7, 1, 8, 3], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 8, 1, 9, 2, 7, 1, 8, and 3.

e 59, step 148000, loss 1.7669554948806763

Target: tensor([9, 1, 9, 9, 1, 0, 0, 8], device='cuda:0')

Pred: tensor([9, 1, 9, 9, 1, 0, 0, 8], device='cuda:0')

9 1 9 9 1 0 0 8

e 60, step 150000, loss 1.7614527940750122

Target: tensor([6, 9, 7, 8, 1, 9, 1, 3], device='cuda:0')

Pred: tensor([6, 9, 7, 8, 1, 9, 1, 3], device='cuda:0')

6 9 7 8 1 9 1 3

e 60, step 152000, loss 1.7723456621170044

Target: tensor([2, 2, 3, 7, 7, 4, 4, 5], device='cuda:0')

Pred: tensor([2, 2, 3, 7, 7, 4, 4, 5], device='cuda:0')

2 2 3 7 7 4 4 5

e 61, step 154000, loss 1.8009408712387085

Target: tensor([2, 3, 9, 7, 6, 7, 4, 5], device='cuda:0')

Pred: tensor([2, 3, 9, 7, 6, 7, 4, 5], device='cuda:0')

2 3 9 7 6 7 4 5

e 62, step 156000, loss 1.8101850748062134

Target: tensor([2, 3, 8, 3, 7, 6, 8, 4], device='cuda:0')

Pred: tensor([2, 3, 8, 3, 7, 6, 8, 4], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 2, 3, 8, 3, 7, 6, 8, and 4.

e 63, step 158000, loss 1.8040566444396973

Target: tensor([7, 6, 8, 3, 3, 7, 6, 3], device='cuda:0')

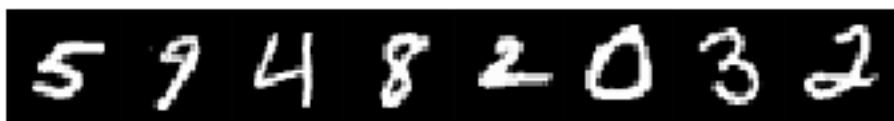
Pred: tensor([7, 6, 8, 3, 3, 7, 6, 3], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 7, 6, 8, 3, 3, 7, 6, and 3.

e 64, step 160000, loss 1.8481130599975586

Target: tensor([5, 9, 4, 8, 2, 0, 3, 2], device='cuda:0')

Pred: tensor([5, 7, 4, 8, 2, 0, 3, 2], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 5, 9, 4, 8, 2, 0, 3, and 2.

e 64, step 162000, loss 1.802116870880127

Target: tensor([5, 7, 1, 5, 7, 6, 2, 7], device='cuda:0')

Pred: tensor([5, 7, 1, 5, 7, 6, 2, 7], device='cuda:0')

A horizontal strip showing eight handwritten digits in white on a black background. The digits are 5, 7, 1, 5, 1, 6, 2, and 7.

e 65, step 164000, loss 1.7994961738586426

Target: tensor([7, 5, 6, 9, 4, 5, 5, 5], device='cuda:0')

Pred: tensor([7, 5, 6, 9, 4, 5, 5, 5], device='cuda:0')

7 5 6 9 4 5 5 5

e 66, step 166000, loss 1.821411371231079

Target: tensor([8, 0, 9, 2, 9, 5, 9, 6], device='cuda:0')

Pred: tensor([8, 0, 4, 2, 9, 5, 9, 6], device='cuda:0')

8 0 9 2 9 5 9 6

e 67, step 168000, loss 1.8479673862457275

Target: tensor([0, 7, 1, 1, 3, 6, 0, 0], device='cuda:0')

Pred: tensor([0, 7, 1, 1, 3, 6, 0, 0], device='cuda:0')

0 7 1 1 3 6 0 0

e 68, step 170000, loss 1.9464595317840576

Target: tensor([0, 7, 8, 8, 7, 3, 1, 8], device='cuda:0')

Pred: tensor([0, 2, 8, 6, 7, 3, 1, 8], device='cuda:0')

0 7 8 6 7 3 1 8

e 68, step 172000, loss 1.8212069272994995

Target: tensor([3, 4, 4, 1, 2, 6, 9, 6], device='cuda:0')

Pred: tensor([3, 4, 4, 1, 2, 6, 9, 6], device='cuda:0')

3 4 4 1 2 6 9 6

e 69, step 174000, loss 1.8676215410232544

Target: tensor([5, 7, 1, 0, 2, 5, 6, 3], device='cuda:0')

Pred: tensor([5, 7, 1, 0, 2, 5, 6, 3], device='cuda:0')

5 7 1 0 2 5 6 3

e 70, step 176000, loss 1.738046646118164

Target: tensor([6, 4, 4, 8, 3, 1, 9, 7], device='cuda:0')

Pred: tensor([6, 4, 4, 8, 3, 1, 9, 7], device='cuda:0')

6 4 4 8 3 1 9 7

e 71, step 178000, loss 1.846350908279419

Target: tensor([0, 3, 7, 3, 2, 3, 2, 5], device='cuda:0')

Pred: tensor([0, 3, 7, 3, 2, 3, 2, 5], device='cuda:0')

0 3 7 3 2 3 2 5

e 72, step 180000, loss 1.8458359241485596

Target: tensor([6, 8, 2, 8, 5, 1, 6, 0], device='cuda:0')

Pred: tensor([6, 8, 2, 8, 5, 1, 6, 0], device='cuda:0')

6 8 2 8 5 1 6 0

e 72, step 182000, loss 1.8131740093231201

Target: tensor([8, 8, 2, 3, 7, 3, 7, 2], device='cuda:0')

Pred: tensor([8, 8, 2, 3, 7, 3, 7, 2], device='cuda:0')

8 8 2 3 7 3 7 2

e 73, step 184000, loss 1.7366836071014404

Target: tensor([4, 9, 2, 7, 4, 9, 6, 6], device='cuda:0')

Pred: tensor([4, 9, 2, 7, 4, 9, 6, 6], device='cuda:0')

4 9 2 7 4 9 6 6

e 74, step 186000, loss 1.7867136001586914

Target: tensor([5, 6, 5, 8, 7, 6, 4, 2], device='cuda:0')

Pred: tensor([5, 6, 5, 8, 7, 6, 4, 2], device='cuda:0')

5 6 5 8 7 6 4 2

e 75, step 188000, loss 1.8313355445861816

Target: tensor([2, 7, 5, 8, 4, 8, 9, 1], device='cuda:0')

Pred: tensor([2, 7, 5, 8, 4, 8, 9, 1], device='cuda:0')

2 7 5 8 4 6 9 1

e 76, step 190000, loss 1.7501091957092285

Target: tensor([7, 0, 0, 3, 5, 6, 9, 6], device='cuda:0')

Pred: tensor([7, 0, 0, 3, 5, 6, 9, 6], device='cuda:0')

7 0 0 3 5 6 9 6

e 76, step 192000, loss 1.9507859945297241

Target: tensor([5, 2, 8, 7, 4, 9, 4, 4], device='cuda:0')

Pred: tensor([5, 2, 8, 7, 4, 9, 4, 4], device='cuda:0')

5 2 8 7 4 9 4 4

e 77, step 194000, loss 1.7771430015563965

Target: tensor([6, 1, 4, 1, 8, 9, 2, 8], device='cuda:0')

Pred: tensor([6, 1, 4, 1, 8, 9, 2, 8], device='cuda:0')

6 1 4 1 8 9 2 8

e 78, step 196000, loss 1.8054965734481812

Target: tensor([8, 4, 0, 1, 2, 4, 5, 7], device='cuda:0')

Pred: tensor([8, 4, 0, 1, 8, 4, 5, 7], device='cuda:0')

8 4 0 1 2 4 5 7

e 79, step 198000, loss 1.8468605279922485

Target: tensor([4, 9, 0, 5, 9, 0, 7, 3], device='cuda:0')

Pred: tensor([4, 7, 0, 5, 9, 0, 7, 3], device='cuda:0')

4 9 0 5 9 0 7 3

e 80, step 200000, loss 1.8777968883514404

Target: tensor([9, 9, 5, 9, 2, 0, 4, 0], device='cuda:0')

Pred: tensor([9, 9, 8, 9, 2, 0, 4, 0], device='cuda:0')

9 9 5 9 2 0 4 0

e 80, step 202000, loss 1.788085699081421

Target: tensor([6, 2, 1, 1, 0, 2, 7, 9], device='cuda:0')

Pred: tensor([6, 2, 1, 1, 0, 2, 7, 9], device='cuda:0')

6 2 1 1 0 2 7 9

e 81, step 204000, loss 1.7849531173706055

Target: tensor([8, 1, 5, 4, 7, 6, 1, 6], device='cuda:0')

Pred: tensor([8, 1, 5, 4, 7, 6, 1, 6], device='cuda:0')

8 1 5 4 7 6 1 6

e 82, step 206000, loss 1.7800381183624268

Target: tensor([3, 6, 7, 4, 1, 1, 0, 5], device='cuda:0')

Pred: tensor([3, 6, 7, 4, 1, 1, 0, 5], device='cuda:0')

3 6 7 4 1 1 0 5

e 83, step 208000, loss 1.7869247198104858

Target: tensor([0, 7, 9, 8, 0, 5, 5, 7], device='cuda:0')

Pred: tensor([0, 7, 9, 8, 0, 5, 5, 7], device='cuda:0')

0 7 9 8 0 5 5 7

e 84, step 210000, loss 1.8137942552566528

Target: tensor([3, 8, 5, 6, 1, 7, 6, 6], device='cuda:0')

Pred: tensor([3, 8, 5, 6, 1, 7, 6, 6], device='cuda:0')

3 8 5 6 1 7 6 6

e 84, step 212000, loss 1.8026354312896729

Target: tensor([3, 7, 1, 6, 5, 8, 1, 5], device='cuda:0')

Pred: tensor([3, 7, 1, 6, 5, 8, 1, 5], device='cuda:0')

3 7 1 6 5 8 1 5

e 85, step 214000, loss 1.8075535297393799

Target: tensor([5, 1, 6, 4, 8, 0, 1, 2], device='cuda:0')

Pred: tensor([5, 1, 6, 4, 8, 0, 1, 2], device='cuda:0')

5 1 6 4 8 0 1 2

e 86, step 216000, loss 1.8540658950805664

Target: tensor([8, 4, 1, 7, 7, 6, 4, 1], device='cuda:0')

Pred: tensor([8, 9, 1, 7, 7, 6, 4, 1], device='cuda:0')

8 4 1 7 7 6 4 1

e 87, step 218000, loss 1.8800668716430664

Target: tensor([6, 6, 0, 3, 8, 5, 2, 2], device='cuda:0')

Pred: tensor([6, 6, 0, 3, 8, 5, 2, 2], device='cuda:0')

6 6 0 3 8 5 2 2

e 88, step 220000, loss 1.7870948314666748

Target: tensor([9, 7, 3, 5, 8, 7, 7, 0], device='cuda:0')

Pred: tensor([9, 7, 3, 5, 8, 7, 7, 0], device='cuda:0')

9 7 3 5 8 7 7 0

e 88, step 222000, loss 1.8001115322113037

Target: tensor([8, 9, 0, 4, 1, 1, 0, 3], device='cuda:0')

Pred: tensor([8, 9, 0, 4, 1, 1, 0, 3], device='cuda:0')

8 9 0 4 1 1 0 3

e 89, step 224000, loss 1.7815284729003906

Target: tensor([2, 5, 9, 4, 4, 0, 9, 3], device='cuda:0')

Pred: tensor([2, 5, 9, 4, 4, 0, 9, 3], device='cuda:0')

2 5 9 4 4 0 9 3

e 90, step 226000, loss 1.8692057132720947

Target: tensor([7, 9, 6, 5, 4, 2, 7, 1], device='cuda:0')

Pred: tensor([7, 9, 6, 5, 4, 2, 7, 1], device='cuda:0')

7 9 6 5 4 2 7 1

e 91, step 228000, loss 1.7200239896774292

Target: tensor([9, 7, 7, 1, 8, 6, 7, 4], device='cuda:0')

Pred: tensor([9, 7, 7, 1, 8, 6, 7, 4], device='cuda:0')

9 7 7 / 8 6 7 4

e 92, step 230000, loss 1.792433738708496
Target: tensor([6, 9, 5, 5, 9, 8, 7, 2], device='cuda:0')
Pred: tensor([6, 9, 5, 5, 9, 8, 7, 2], device='cuda:0')

6 9 5 5 9 8 7 2

e 92, step 232000, loss 1.7548162937164307
Target: tensor([0, 8, 4, 1, 0, 1, 1, 7], device='cuda:0')
Pred: tensor([0, 8, 4, 1, 0, 1, 1, 7], device='cuda:0')

0 8 4 1 0 1 1 7

e 93, step 234000, loss 1.8704413175582886
Target: tensor([0, 4, 8, 7, 1, 3, 2, 0], device='cuda:0')
Pred: tensor([0, 4, 8, 7, 2, 3, 2, 0], device='cuda:0')

0 4 8 7 1 3 2 0

e 94, step 236000, loss 1.793448805809021
Target: tensor([0, 4, 2, 3, 1, 3, 1, 9], device='cuda:0')
Pred: tensor([0, 4, 2, 3, 1, 3, 1, 9], device='cuda:0')

0 4 2 3 1 3 1 9

e 95, step 238000, loss 1.7790563106536865

Target: tensor([1, 3, 5, 3, 1, 3, 0, 5], device='cuda:0')

Pred: tensor([1, 3, 5, 3, 1, 3, 0, 5], device='cuda:0')

1 3 5 3 1 3 0 5

e 96, step 240000, loss 1.8358275890350342

Target: tensor([6, 6, 9, 6, 7, 6, 6, 0], device='cuda:0')

Pred: tensor([6, 6, 9, 6, 7, 6, 6, 0], device='cuda:0')

6 6 9 6 7 6 6 0

e 96, step 242000, loss 1.8352785110473633

Target: tensor([4, 5, 3, 4, 3, 3, 6, 0], device='cuda:0')

Pred: tensor([4, 5, 3, 4, 3, 3, 6, 0], device='cuda:0')

4 5 3 4 3 3 6 0

e 97, step 244000, loss 1.8189873695373535

Target: tensor([9, 8, 6, 4, 8, 8, 3, 7], device='cuda:0')

Pred: tensor([9, 8, 6, 4, 8, 8, 3, 7], device='cuda:0')

9 8 6 4 8 8 3 7

e 98, step 246000, loss 1.8110897541046143

Target: tensor([2, 6, 3, 9, 8, 9, 9, 1], device='cuda:0')

Pred: tensor([2, 6, 3, 9, 8, 9, 9, 1], device='cuda:0')

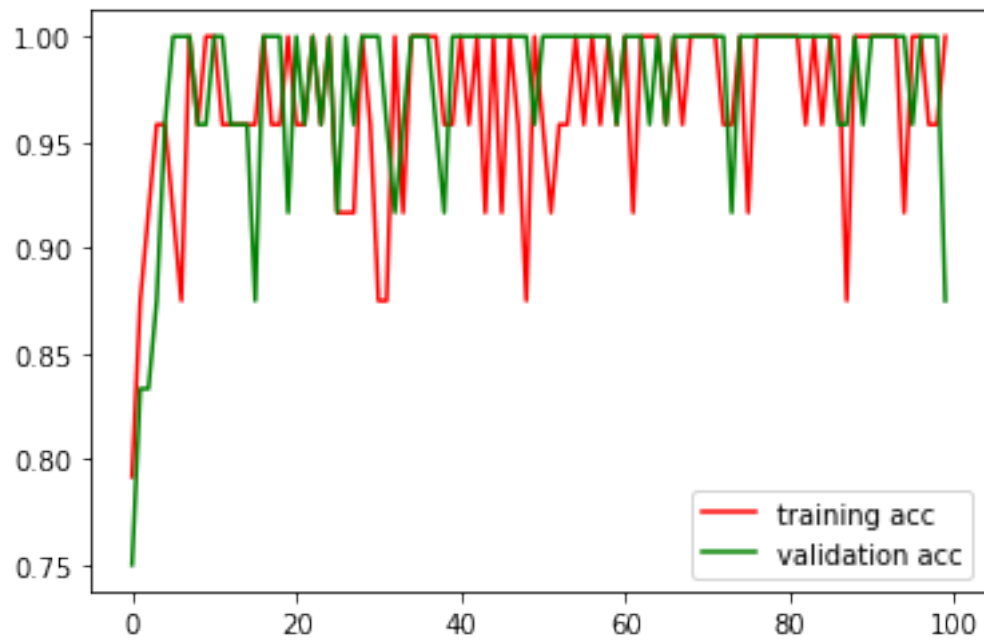

2 6 3 9 8 9 9 1

e 99, step 248000, loss 1.7869174480438232

Target: tensor([4, 0, 2, 5, 4, 2, 2, 2], device='cuda:0')

Pred: tensor([4, 0, 2, 5, 4, 2, 2, 2], device='cuda:0')

4 0 2 5 4 2 2 2



1.15 Submission

- Make sure you have finished all the required implementation tasks.
- Check your codes and make sure the result in each section could be reproducible.

- Upload the Jupyter file with all required figures plotted.
- Upload a pdf version of this Jupyter note. You can first download a html file by clicking on the Jupyter menu bar: File -> Download as -> HTML (.html). Then open the html file and convert it into a pdf file with your browser.

[]: